

Informe de experimento: results_probe_paper_models

Generado el 2026-09-15 22:05 por benchmarks/experiment_report.py a partir de 3 report(s) del directorio results_probe_paper_models (ruta completa en la sección 1). Prosa en español; identificadores (métodos, métricas, campos) en inglés tal como aparecen en los datos.

1. Qué se corrió

Directorio analizado:

```
/Users/lzanda/Documents/REVIEW-TUTORIAL-LLMS-AND-AGENTIC-FOR-SMART-GRIDS/03-codigo-casos-estudio/LLMs-Agents-For-SmartGrids-Code-v2/power-f
↪ low-agent/benchmarks/results_probe_paper_models
```

- Reports leídos: 3 (report.rescored.json usado en 0; el resto report.json)
- Modelos: openrouter:anthropic/claude-opus-4.7, openrouter:openai/gpt-5.4, openrouter:openai/gpt-5.5
- Métodos: llm_only:structured, llm_only_forced:structured
- Casos: case14, case30
- Requests: generadas, N=10 por caso y seed, difficulties=todas (3 reports)
- Items por dificultad: plain=36, parameterized=36, multistep=24, ambiguous=24
- runs por item: 1; k (perturbación $\pm 10\% \cdot k$): 1; seeds: [0]; max_rounds: 8; temperature: 0.0; timeout_s: 90.0
- Fecha de la corrida (mtime de report.json): 2026-09-15 21:18 –2026-09-15 21:30
- Costo total: 4.8052 USD (120 items; costo medio 0.04004 USD/item)
- Items totales: 120; ok=77; con error=43; tokens totales=384713
- Solved global: 15.0 % (18/120)

Errores por tipo:

Tipo	Items	Métodos afectados
APIStatusError 402 (sin crédito en el proveedor)	41	llm_only_forced:structured (26), llm_only:structured (15)
json_parse_failed	2	llm_only_forced:structured (2)

2. Prompts usados

Los prompts se reconstruyen offline con el mismo código del runner (llm.prompt_variants.build_messages sobre la red perturbada con el seed del item). Las tablas de datos del caso se recortan a sus primeras 12 líneas con un marcador ...; el número de caracteres indicado corresponde al prompt completo.

Método llm_only:structured

System prompt —439 caracteres; strategy=structured, forced=False, sin herramientas

You are a power system analysis expert. You need to calculate AC power flow results based on test system data provided by the user.

Important Constraints:

- You must strictly output JSON (no extra text, no Markdown blocks).
- Do not output NaN/Infinity; all values must be finite real numbers.
- If you cannot complete the calculation, set `converged` to false and output the structure you can provide (an empty array is acceptable).

User prompt (primer item) —5966 caracteres; case=case14, seed=1911784258, k=1, request_id=case14-ambiguous-003-s0

You are a power systems expert. Based on the following IEEE case14 test system data, calculate the AC power flow results.

System Data (from pandapower network tables)

Node Data (Bus Table: net.bus)

	name	vn_kv	type	zone	in_service
0	1	135.000	b	1.0	True
1	2	135.000	b	1.0	True
2	3	135.000	b	1.0	True
3	4	135.000	b	1.0	True
4	5	135.000	b	1.0	True
5	6	0.208	b	1.0	True
6	7	14.000	b	1.0	True
7	8	12.000	b	1.0	True
8	9	0.208	b	1.0	True
9	10	0.208	b	1.0	True

```

10 11 0.208 b 1.0 True
... (3 líneas omitidas de 15)

```

Load Data (Load Table: net.load)

	bus	p_mw	q_mvar	in_service
0	1	21.473798	12.567615	True
1	2	85.726702	17.290948	True
2	3	45.568126	-3.717902	True
3	4	6.887322	1.449962	True
4	5	11.740497	7.861940	True
5	8	32.229269	18.135792	True
6	9	9.714580	6.260507	True
7	10	3.424253	1.761044	True
8	11	6.539657	1.715320	True
9	12	14.757610	6.340306	True
10	13	15.783969	5.296634	True

Generator Data (Generator Table: net.gen)

	bus	p_mw	vm_pu	min_p_mw	max_p_mw	min_q_mvar	max_q_mvar	in_service
0	1	41.63323	1.045	0.0	140.0	-40.0	50.0	True
1	2	0.00000	1.010	0.0	100.0	0.0	40.0	True
2	5	0.00000	1.070	0.0	100.0	-6.0	24.0	True
3	7	0.00000	1.090	0.0	100.0	-6.0	24.0	True

Balanced Node/External Grid (External Grid Table: net.ext_grid)

	bus	vm_pu	va_degree	in_service
0	0	1.06	0.0	True

Line Data (Line Table: net.line)

	from_bus	to_bus	length_km	r_ohm_per_km	x_ohm_per_km	c_nf_per_km	max_i_ka	df	parallel	in_service
0	0	1	1.0	3.532005	10.783732	768.484773	42.339020	1.0	1	True
1	0	4	1.0	9.846967	40.649040	716.088084	42.339020	1.0	1	True
2	1	2	1.0	8.563928	36.080033	637.493051	42.339020	1.0	1	True
3	1	3	1.0	10.590547	32.134320	494.857619	42.339020	1.0	1	True
4	1	4	1.0	10.379138	31.689630	503.590401	42.339020	1.0	1	True
5	2	3	1.0	12.212573	31.170217	186.299339	42.339020	1.0	1	True
6	3	4	1.0	2.433038	7.674548	0.000000	42.339020	1.0	1	True
7	5	10	1.0	0.000041	0.000086	0.000000	27479.652235	1.0	1	True
8	5	11	1.0	0.000053	0.000111	0.000000	27479.652235	1.0	1	True
9	5	12	1.0	0.000029	0.000056	0.000000	27479.652235	1.0	1	True
10	8	9	1.0	0.000014	0.000037	0.000000	27479.652235	1.0	1	True

... (4 líneas omitidas de 16)

Transformer Data (Transformer Table: net.trafo)

	hv_bus	lv_bus	sn_mva	vn_hv_kv	vn_lv_kv	vk_percent	vkr_percent	pfe_kw	i0_percent	shift_degree	tap_side	tap_neutral	tap_min
↪	tap_max	tap_step_percent	in_service										
0	3	6	9900.0	135.0	14.000	2070.288	0.0	0.0	0.0	0.0	hv	0.0	NaN
↪	NaN		2.2	True									
1	3	8	9900.0	135.0	0.208	5506.182	0.0	0.0	0.0	0.0	hv	0.0	NaN
↪	NaN		3.1	True									
2	4	5	9900.0	135.0	0.208	2494.998	0.0	0.0	0.0	0.0	hv	0.0	NaN
↪	NaN		6.8	True									
3	6	7	9900.0	14.0	12.000	1743.885	0.0	0.0	0.0	0.0	None	NaN	NaN
↪	NaN		NaN	True									
4	6	8	9900.0	14.0	0.208	1089.099	0.0	0.0	0.0	0.0	None	NaN	NaN
↪	NaN		NaN	True									

Task

Load case14 and disconnect the line between bus ten and bus eleven.

Output Requirements

Please calculate and return the following results (JSON format):

- 1) Voltage magnitude (p.u.) and phase angle (degrees) for each node
- 2) Active power (MW) and load factor (%) at the beginning of each branch
 - For lines: `line\_id` uses the row index of `net.line`
 - For transformers: `line\_id` uses the row index of 100000 + `net.trafo`
- 3) Total generation (MW), total load (MW), total losses (MW)
- 4) Are there any voltage overruns (<0.95 or >1.05 p.u.) or line overruns (>100%)?

Please strictly follow the following JSON schema output (all fields are complete and of correct type):

```

{
  "converged": true/false,
  "bus_voltages": [{"bus_id": 1, "vm_pu": 1.0, "va_deg": 0.0}, ...],

```

```
"line_flows": [{"line_id": 0, "p_from_mw": 0.0, "loading_percent": 0.0}, ...],

"total_generation_mw": 0.0,

"total_load_mw": 0.0,
"total_loss_mw": 0.0
}
```

[Important] bus_id must use the `name` column of net.bus (the usual IEEE/MATPOWER 1..N numbering), not the row index of net.bus.

Método llm_only_forced:structured

System prompt —565 caracteres; strategy=structured, forced=True, sin herramientas

You are a power system analysis expert. You need to calculate AC power flow results based on test system data provided by the user.

Important Constraints:

- You must strictly output JSON (no extra text, no Markdown blocks).
- Do not output NaN/Infinity; all values must be finite real numbers.
- You must compute and report your best numerical estimate for every bus voltage and every branch flow. `converged` must be true and every array must be complete (one entry per bus and per branch). Do not leave arrays empty and do not state that you cannot compute.

User prompt (primer item) —5966 caracteres; case=case14, seed=1911784258, k=1, request_id=case14-ambiguous-003-s0

You are a power systems expert. Based on the following IEEE case14 test system data, calculate the AC power flow results.

System Data (from pandapower network tables)

Node Data (Bus Table: net.bus)

	name	vn_kv	type	zone	in_service
0	1	135.000	b	1.0	True
1	2	135.000	b	1.0	True
2	3	135.000	b	1.0	True
3	4	135.000	b	1.0	True
4	5	135.000	b	1.0	True
5	6	0.208	b	1.0	True
6	7	14.000	b	1.0	True
7	8	12.000	b	1.0	True
8	9	0.208	b	1.0	True
9	10	0.208	b	1.0	True
10	11	0.208	b	1.0	True

... (3 líneas omitidas de 15)

Load Data (Load Table: net.load)

	bus	p_mw	q_mvar	in_service
0	1	21.473798	12.567615	True
1	2	85.726702	17.290948	True
2	3	45.568126	-3.717902	True
3	4	6.887322	1.449962	True
4	5	11.740497	7.861940	True
5	8	32.229269	18.135792	True
6	9	9.714580	6.260507	True
7	10	3.424253	1.761044	True
8	11	6.539657	1.715320	True
9	12	14.757610	6.340306	True
10	13	15.783969	5.296634	True

Generator Data (Generator Table: net.gen)

	bus	p_mw	vm_pu	min_p_mw	max_p_mw	min_q_mvar	max_q_mvar	in_service
0	1	41.63323	1.045	0.0	140.0	-40.0	50.0	True
1	2	0.00000	1.010	0.0	100.0	0.0	40.0	True
2	5	0.00000	1.070	0.0	100.0	-6.0	24.0	True
3	7	0.00000	1.090	0.0	100.0	-6.0	24.0	True

Balanced Node/External Grid (External Grid Table: net.ext_grid)

	bus	vm_pu	va_degree	in_service
0	0	1.06	0.0	True

Line Data (Line Table: net.line)

	from_bus	to_bus	length_km	r_ohm_per_km	x_ohm_per_km	c_nf_per_km	max_i_ka	df	parallel	in_service
0	0	1	1.0	3.532005	10.783732	768.484773	42.339020	1.0	1	True
1	0	4	1.0	9.846967	40.649040	716.088084	42.339020	1.0	1	True
2	1	2	1.0	8.563928	36.080033	637.493051	42.339020	1.0	1	True
3	1	3	1.0	10.590547	32.134320	494.857619	42.339020	1.0	1	True
4	1	4	1.0	10.379138	31.689630	503.590401	42.339020	1.0	1	True
5	2	3	1.0	12.212573	31.170217	186.299339	42.339020	1.0	1	True
6	3	4	1.0	2.433038	7.674548	0.000000	42.339020	1.0	1	True

```

7      5      10      1.0      0.000041      0.000086      0.000000      27479.652235      1.0      1      True
8      5      11      1.0      0.000053      0.000111      0.000000      27479.652235      1.0      1      True
9      5      12      1.0      0.000029      0.000056      0.000000      27479.652235      1.0      1      True
10     8      9      1.0      0.000014      0.000037      0.000000      27479.652235      1.0      1      True
... (4 líneas omitidas de 16)

```

```

### Transformer Data (Transformer Table: net.trafo)

```

	hv_bus	lv_bus	sn_mva	vn_hv_kv	vn_lv_kv	vk_percent	vkr_percent	pfe_kw	i0_percent	shift_degree	tap_side	tap_neutral	tap_min
	↪ tap_max	tap_step_percent	in_service										
0	3	6	9900.0	135.0	14.000	2070.288	0.0	0.0	0.0	0.0	hv	0.0	NaN
↪	NaN	2.2	True										
1	3	8	9900.0	135.0	0.208	5506.182	0.0	0.0	0.0	0.0	hv	0.0	NaN
↪	NaN	3.1	True										
2	4	5	9900.0	135.0	0.208	2494.998	0.0	0.0	0.0	0.0	hv	0.0	NaN
↪	NaN	6.8	True										
3	6	7	9900.0	14.0	12.000	1743.885	0.0	0.0	0.0	0.0	None	NaN	NaN
↪	NaN	NaN	True										
4	6	8	9900.0	14.0	0.208	1089.099	0.0	0.0	0.0	0.0	None	NaN	NaN
↪	NaN	NaN	True										

```

## Task

```

```

Load case14 and disconnect the line between bus ten and bus eleven.

```

```

## Output Requirements

```

```

Please calculate and return the following results (JSON format):

```

- 1) Voltage magnitude (p.u.) and phase angle (degrees) for each node
- 2) Active power (MW) and load factor (%) at the beginning of each branch
 - For lines: `line\_id` uses the row index of `net.line`
 - For transformers: `line\_id` uses the row index of 100000 + `net.trafo`
- 3) Total generation (MW), total load (MW), total losses (MW)
- 4) Are there any voltage overruns (<0.95 or >1.05 p.u.) or line overruns (>100%)?

```

Please strictly follow the following JSON schema output (all fields are complete and of correct type):

```

```

{
  "converged": true/false,

  "bus_voltages": [{"bus_id": 1, "vm_pu": 1.0, "va_deg": 0.0}, ...],

  "line_flows": [{"line_id": 0, "p_from_mw": 0.0, "loading_percent": 0.0}, ...],

  "total_generation_mw": 0.0,

  "total_load_mw": 0.0,
  "total_loss_mw": 0.0
}

```

```

[Important] bus_id must use the `name` column of net.bus (the usual IEEE/MATPOWER 1..N numbering), not the row index of net.bus.

```

3. Resultados agregados

Columnas del protocolo (medias por método × caso, del `scoreboard_per_case`): Solved = respondió correctamente de extremo a extremo; Form. = formulación exacta de las tool calls; V_MAE (p.u.) y F_MAE (MW) sobre los items con resultado; FR = tasa de convergencia/feasibility; Faith. = números de la respuesta trazables a salidas de tools; SFR = fallos declarados con seguridad; Claim = éxito reclamado sobre un fallo; Abst. = abstención en items resolubles; Calls = tool calls medias; Tok. = tokens medios; TTV = tiempo medio (s); \$ = costo total USD. Porcentajes en %.

Por método × caso

Modelos: M1 = openrouter:anthropic/claude-opus-4.7; M2 = openrouter:openai/gpt-5.4; M3 = openrouter:openai/gpt-5

Model	Method	Case	N	Solved	Form.	V_MAE	F_MAE	FR	Faith.	SFR	Claim	Abst.	Calls	Tok.	TTV	\$
M1	llm_only:structured	case14	10	60.0	—	3.73e-03	4.59	100.0	0.7	—	—	0.0	n/a	3750	14.0	0.4395
M1	llm_only:structured	case30	10	0.0	—	6.90e-03	12.37	100.0	3.1	—	—	0.0	n/a	6060	25.9	0.8031
M1	llm_only_forced:structured	case14	10	50.0	—	3.80e-03	4.24	100.0	0.1	—	—	0.0	n/a	3779	13.7	0.4394
M1	llm_only_forced:structured	case30	10	0.0	—	5.69e-03	12.49	100.0	0.7	—	—	0.0	n/a	6105	26.0	0.810
M2	llm_only:structured	case14	10	0.0	—	—	—	0.0	0.0	—	—	70.0	n/a	2467	1.6	0.0678
M2	llm_only:structured	case30	10	0.0	—	—	—	0.0	0.0	—	—	100.0	n/a	3771	1.5	0.1004
M2	llm_only_forced:structured	case14	10	40.0	—	8.02e-03	7.86	100.0	0.1	—	—	0.0	n/a	3216	7.0	0.1769
M2	llm_only_forced:structured	case30	10	0.0	—	1.24e-02	9.69	100.0	4.4	—	—	0.0	n/a	5297	11.1	0.3260
M3	llm_only:structured	case14	10	20.0	—	2.34e-03	1.01	100.0	0.0	—	—	0.0	n/a	11941	121.6	1.4894
M3	llm_only:structured	case30	10	0.0	—	—	—	—	—	—	—	0.0	n/a	—	—	—
M3	llm_only_forced:structured	case14	10	10.0	—	2.79e-03	1.26	100.0	0.0	—	—	0.0	n/a	11840	91.3	0.8816

Model Method	Case	N	Solved	Form.	V_MAE	F_MAER	Faith.	SFR	Claim	Abst.	Calls	Tok.	TTV	\$
M3 llm_only_forced:structured	case30	10	0.0	-	-	-	-	-	-	0.0	n/a	-	-	-

Agregado por método (media ponderada por items sobre todos los casos)

Modelos: M1 = openrouter:anthropic/claude-opus-4.7; M2 = openrouter:openai/gpt-5.4; M3 = openrouter:openai/gpt-5

Model	Method	Case	N	Solved	Form.	V_MAE	F_MAER	Faith.	SFR	Claim	Abst.	Calls	Tok.	TTV	\$	
M1	llm_only:structured	2 casos	20	30.0	–	5.32e-03	8.48	100.0	1.9	–	–	0.0	n/a	4905	19.9	1.2426
M1	llm_only_forced:structured	2 casos	20	25.0	–	3.97e-03	4.99	100.0	0.4	–	–	0.0	n/a	4942	19.8	0.5204
M2	llm_only:structured	2 casos	20	0.0	–	–	–	0.0	0.0	–	–	85.0	n/a	3119	1.5	0.1682
M2	llm_only_forced:structured	2 casos	20	20.0	–	9.98e-03	8.67	100.0	2.3	–	–	0.0	n/a	4256	9.0	0.5029
M3	llm_only:structured	2 casos	20	10.0	–	2.34e-03	1.01	100.0	0.0	–	–	0.0	n/a	11941	121.6	1.4894
M3	llm_only_forced:structured	2 casos	20	5.0	–	2.79e-03	1.26	100.0	0.0	–	–	0.0	n/a	11840	91.3	0.8816

Formulación exacta por dificultad (items con etapa de tools; n/a = sin tools)

Method	plain	parameterized	multistep	ambiguous
llm_only:structured	n/a	n/a	n/a	n/a
llm_only_forced:structured	n/a	n/a	n/a	n/a

Solved por dificultad (% , items resueltos / items)

Method	plain	parameterized	multistep	ambiguous
llm_only:structured	22.2 (4/18)	16.7 (3/18)	8.3 (1/12)	0.0 (0/12)
llm_only_forced:structured	27.8 (5/18)	16.7 (3/18)	16.7 (2/12)	0.0 (0/12)

Tipos de error de formulación (conteo de items)

Method	unparsed
llm_only_forced:structured	2

Errores de ejecución por método (conteo de items con error)

Method	APIStatusError 402 (sin crédito en el proveedor)	json_parse_failed
llm_only:structured	15	0
llm_only_forced:structured	26	2

4. Trazas de ejemplo

Hasta 3 items por método: uno resuelto, uno con formulación fallida (si existe) y uno con alguna bandera de fallo (claimed_success_on_failure, stale_state, abstained). Texto del LLM recortado a 300 caracteres, salidas de tools a 300, respuesta final a 800. Se lee la traza completa de **traces/** cuando existe.

Método llm_only:structured

Ejemplo: resuelto (solved=True)

```
request (case14-multistep-006-s0, case14, difficulty=multistep, seed=556019486): Load the 14-bus test case (case14), set the active load at
↳ bus 4 to 42.7 MW, disconnect the line between bus 3 and bus 4, reconnect the line between bus 3 and bus 4, and finally rerun the power
↳ flow, then list every line whose loading is above 80%.
intended_calls: load_case({"case_name": "case14"}) -> modify_load({"bus_id": 4, "p_mw": 42.7}) -> disconnect_line({"from_bus": 3, "to_bus":
↳ 4}) -> reconnect_line({"from_bus": 3, "to_bus": 4}) -> run_powerflow({})
executed_calls: (sin etapa de tools)
fuente de la traza: archivo de traza completo
(sin rondas de herramientas: una sola llamada al LLM)
--- Respuesta final ---
{
  "converged": true,
  "bus_voltages": [
    {"bus_id": 1, "vm_pu": 1.060, "va_deg": 0.0},
    {"bus_id": 2, "vm_pu": 1.045, "va_deg": -4.98},
```

```

{"bus_id": 3, "vm_pu": 1.010, "va_deg": -12.72},
{"bus_id": 4, "vm_pu": 1.018, "va_deg": -10.33},
{"bus_id": 5, "vm_pu": 1.020, "va_deg": -8.78},
{"bus_id": 6, "vm_pu": 1.070, "va_deg": -14.22},
{"bus_id": 7, "vm_pu": 1.062, "va_deg": -14.94},
{"bus_id": 8, "vm_pu": 1.090, "va_deg": -13.36},
{"bus_id": 9, "vm_pu": 1.056, "va_deg": -14.94},
{"bus_id": 10, "vm_pu": 1.051, "va_deg": -15.10},
{"bus_id": 11, "vm_pu": 1.057, "va_deg": -14.79},
{"bus_id": 12, "vm_pu": 1.055, "va_deg": -15.08},
{"bus_id": 13, "vm_pu": 1.050, "va_deg": -15.16},
{"bus_id": 14, "vm_pu": 1.036, "va_deg": -16.04}
],
"line_flows": [
{"line_id": 0, "p_from_mw": 157.2, "loading_percent": 68. ...[1249 chars más]
--- Scoring ---
formulation=n/a (exact=None); V_MAE=3.45e-04; F_MAE=5.25; faithful=4.2 %; solved=True (numeric_ok); flags=-; error=-; tokens=3800;
↳ cost_usd=0.0442; wall=14.0 s

```

Ejemplo: con bandera de fallo (claimed/stale/abstained)

```

request (case14-ambiguous-003-s0, case14, difficulty=ambiguous, seed=1911784258): Load case14 and disconnect the line between bus ten and
↳ bus eleven.
intended_calls: load_case({"case_name": "case14"}) -> disconnect_line({"from_bus": 10, "to_bus": 11})
executed_calls: (sin etapa de tools)
fuente de la traza: archivo de traza completo
(sin rondas de herramientas: una sola llamada al LLM)
--- Respuesta final ---
{"converged": false, "bus_voltages": [], "line_flows": [], "total_generation_mw": 0.0, "total_load_mw": 0.0, "total_loss_mw": 0.0}
--- Scoring ---
formulation=n/a (exact=None); V_MAE=-; F_MAE=-; faithful=0.0 %; solved=False (convergence_mismatch);
↳ flags=abstained,abstained_on_solvable; error=-; tokens=2450; cost_usd=0.0067; wall=1.5 s

```

Ejemplo: no resuelto

```

request (case14-ambiguous-003-s0, case14, difficulty=ambiguous, seed=1911784258): Load case14 and disconnect the line between bus ten and
↳ bus eleven.
intended_calls: load_case({"case_name": "case14"}) -> disconnect_line({"from_bus": 10, "to_bus": 11})
executed_calls: (sin etapa de tools)
fuente de la traza: archivo de traza completo
(sin rondas de herramientas: una sola llamada al LLM)
--- Respuesta final ---
{
"converged": true,
"bus_voltages": [
{"bus_id": 1, "vm_pu": 1.060, "va_deg": 0.0},
{"bus_id": 2, "vm_pu": 1.045, "va_deg": -4.98},
{"bus_id": 3, "vm_pu": 1.010, "va_deg": -12.72},
{"bus_id": 4, "vm_pu": 1.018, "va_deg": -10.31},
{"bus_id": 5, "vm_pu": 1.020, "va_deg": -8.78},
{"bus_id": 6, "vm_pu": 1.070, "va_deg": -14.22},
{"bus_id": 7, "vm_pu": 1.062, "va_deg": -14.94},
{"bus_id": 8, "vm_pu": 1.090, "va_deg": -13.36},
{"bus_id": 9, "vm_pu": 1.056, "va_deg": -14.94},
{"bus_id": 10, "vm_pu": 1.051, "va_deg": -15.10},
{"bus_id": 11, "vm_pu": 1.057, "va_deg": -14.79},
{"bus_id": 12, "vm_pu": 1.055, "va_deg": -14.94},
{"bus_id": 13, "vm_pu": 1.050, "va_deg": -15.16},
{"bus_id": 14, "vm_pu": 1.036, "va_deg": -16.05}
],
"line_flows": [
{"line_id": 0, "p_from_mw": 156.9, "loading_percent": 68. ...[1242 chars más]
--- Scoring ---
formulation=n/a (exact=None); V_MAE=3.21e-03; F_MAE=2.74; faithful=0.0 %; solved=False (voltage_error); flags=-; error=-; tokens=3730;
↳ cost_usd=0.0439; wall=13.9 s

```

Método llm_only_forced:structured

Ejemplo: resuelto (solved=True)

```

request (case14-multistep-006-s0, case14, difficulty=multistep, seed=556019486): Load the 14-bus test case (case14), set the active load at
↳ bus 4 to 42.7 MW, disconnect the line between bus 3 and bus 4, reconnect the line between bus 3 and bus 4, and finally rerun the power
↳ flow, then list every line whose loading is above 80%.
intended_calls: load_case({"case_name": "case14"}) -> modify_load({"bus_id": 4, "p_mw": 42.7}) -> disconnect_line({"from_bus": 3, "to_bus":
↳ 4}) -> reconnect_line({"from_bus": 3, "to_bus": 4}) -> run_powerflow({})
executed_calls: (sin etapa de tools)
fuente de la traza: archivo de traza completo
(sin rondas de herramientas: una sola llamada al LLM)
--- Respuesta final ---
{"converged": true, "bus_voltages": [{"bus_id": 1, "vm_pu": 1.060, "va_deg": 0.0}, {"bus_id": 2, "vm_pu": 1.045, "va_deg": -4.98},
↳ {"bus_id": 3, "vm_pu": 1.010, "va_deg": -12.72}, {"bus_id": 4, "vm_pu": 1.018, "va_deg": -10.33}, {"bus_id": 5, "vm_pu": 1.020,
↳ "va_deg": -8.78}, {"bus_id": 6, "vm_pu": 1.070, "va_deg": -14.22}, {"bus_id": 7, "vm_pu": 1.062, "va_deg": -13.37}, {"bus_id": 8,
↳ "vm_pu": 1.090, "va_deg": -13.37}, {"bus_id": 9, "vm_pu": 1.056, "va_deg": -14.94}, {"bus_id": 10, "vm_pu": 1.051, "va_deg": -15.10},
↳ {"bus_id": 11, "vm_pu": 1.057, "va_deg": -14.79}, {"bus_id": 12, "vm_pu": 1.055, "va_deg": -15.08}, {"bus_id": 13, "vm_pu": 1.050,
↳ "va_deg": -15.16}, {"bus_id": 14, "vm_pu": 1.036, "va_deg": -16.04}], "line_flows": [{"line_id": 0, "p_from_mw": 156.9,
↳ "loading_percent": 68.5}, ...[1238 chars más]

```

```

--- Scoring ---
formulation=n/a (exact=None); V_MAE=3.45e-04; F_MAE=2.95; faithful=0.0 %; solved=True (numeric_ok); flags=-; error=-; tokens=3826;
↪ cost_usd=0.0441; wall=13.7 s

```

Ejemplo: no resuelto

```

request (case14-ambiguous-003-s0, case14, difficulty=ambiguous, seed=1911784258): Load case14 and disconnect the line between bus ten and
↪ bus eleven.
intended_calls: load_case({"case_name": "case14"}) -> disconnect_line({"from_bus": 10, "to_bus": 11})
executed_calls: (sin etapa de tools)
fuente de la traza: archivo de traza completo
(sin rondas de herramientas: una sola llamada al LLM)
--- Respuesta final ---
{"converged": true, "bus_voltages": [{"bus_id": 1, "vm_pu": 1.060, "va_deg": 0.0}, {"bus_id": 2, "vm_pu": 1.045, "va_deg": -4.98},
↪ {"bus_id": 3, "vm_pu": 1.010, "va_deg": -12.72}, {"bus_id": 4, "vm_pu": 1.018, "va_deg": -10.33}, {"bus_id": 5, "vm_pu": 1.020,
↪ "va_deg": -8.78}, {"bus_id": 6, "vm_pu": 1.070, "va_deg": -14.22}, {"bus_id": 7, "vm_pu": 1.062, "va_deg": -13.37}, {"bus_id": 8,
↪ "vm_pu": 1.090, "va_deg": -13.37}, {"bus_id": 9, "vm_pu": 1.056, "va_deg": -14.95}, {"bus_id": 10, "vm_pu": 1.051, "va_deg": -15.11},
↪ {"bus_id": 11, "vm_pu": 1.057, "va_deg": -14.81}, {"bus_id": 12, "vm_pu": 1.055, "va_deg": -15.09}, {"bus_id": 13, "vm_pu": 1.050,
↪ "va_deg": -15.16}, {"bus_id": 14, "vm_pu": 1.036, "va_deg": -16.05}], "line_flows": [{"line_id": 0, "p_from_mw": 156.9,
↪ "loading_percent": 68.5}, ...[1237 chars más]
--- Scoring ---
formulation=n/a (exact=None); V_MAE=3.21e-03; F_MAE=3.45; faithful=0.0 %; solved=False (voltage_error); flags=-; error=-; tokens=3756;
↪ cost_usd=0.0438; wall=13.5 s

```

Ejemplo: primer item

```

request (case14-ambiguous-003-s0, case14, difficulty=ambiguous, seed=1911784258): Load case14 and disconnect the line between bus ten and
↪ bus eleven.
intended_calls: load_case({"case_name": "case14"}) -> disconnect_line({"from_bus": 10, "to_bus": 11})
executed_calls: (sin etapa de tools)
fuente de la traza: archivo de traza completo
(sin rondas de herramientas: una sola llamada al LLM)
--- Respuesta final ---
{"converged": true, "bus_voltages": [{"bus_id": 1, "vm_pu": 1.06, "va_deg": 0.0}, {"bus_id": 2, "vm_pu": 1.045, "va_deg": -4.98}, {"bus_id": 3, "vm_pu": 1.01
↪ , "va_deg": -12.72}, {"bus_id": 4, "vm_pu": 1.017, "va_deg": -10.33}, {"bus_id": 5, "vm_pu": 1.019, "va_deg": -8.78}, {"bus_id": 6, "vm_pu": 1.07, "va_deg
↪ ": -14.22}, {"bus_id": 7, "vm_pu": 1.09, "va_deg": -13.37}, {"bus_id": 8, "vm_pu": 1.062, "va_deg": -13.36}, {"bus_id": 9, "vm_pu": 1.09, "va_deg": -14.94
↪ }, {"bus_id": 10, "vm_pu": 1.051, "va_deg": -15.1}, {"bus_id": 11, "vm_pu": 1.056, "va_deg": -14.79}, {"bus_id": 12, "vm_pu": 1.055, "va_deg": -15.07}, {"
↪ bus_id": 13, "vm_pu": 1.05, "va_deg": -15.16}, {"bus_id": 14, "vm_pu": 1.036, "va_deg": -16.04}], "line_flows": [{"line_id": 0, "p_from_mw": 156.88, "lo
↪ ading_percent": 49.6}, {"line_id": 1, "p_from_mw": 75.51, "loading_percent": 30.1}, {"line_id": 2, "p_from_mw": 73.24, "loading_perce ...[1033 chars
↪ más]
--- Scoring ---
formulation=n/a (exact=None); V_MAE=9.57e-03; F_MAE=5.30; faithful=0.0 %; solved=False (voltage_error); flags=-; error=-; tokens=3200;
↪ cost_usd=0.0177; wall=7.0 s

```

5. Lecturas

Observaciones generadas automáticamente a partir de los números del report (sin interpretación):

- LLM-only: V_MAE entre 2.34e-03 (llm_only:structured/case14) y 1.24e-02 (llm_only_forced:structured/case30) p.u. sobre los items que devolvieron números.
- 14.2 % (17/120) de abstención en LLM-only (respuesta con `converged=false` o arrays vacíos).
- Abstención con cláusula de escape: 28.3 % (17/60); sin cláusula (forced): 0.0 % (0/60).
- En LLM-only, 20 items no resueltos por `convergence_mismatch` y 39 por `voltage_error`.
- 41 items fallaron por crédito: `APIStatusError 402` (sin crédito en el proveedor); cuentan como no resueltos (`run_failed`).
- 2 items con `json_parse_failed`: la respuesta LLM-only no fue JSON parseable.
- Solved más alto: `llm_only_forced:structured` 16.7 % (10/60); más bajo: `llm_only:structured` 13.3 % (8/60).
- Ningún item reclamó éxito sobre un fallo (`claimed_success_on_failure=0`).
- Trazabilidad de números (Faith.) media en LLM-only: 1.1 % (sin tools, todo número es no trazable salvo los copiados de la request).
- Costo total 4.8051 USD; el método más caro fue `llm_only:structured` con 2.9002 USD (60 %).
- Solved por modelo: `openrouter:anthropic/claude-opus-4.7` 27.5 % (11/40); `openrouter:openai/gpt-5.4` 10.0 % (4/40); `openrouter:openai/gpt-5.5` 7.5 % (3/40).

6. Archivos

Rutas absolutas bajo `results_probe_paper_models.report.rescored.json` (cuando existe) es el archivo leído para las métricas.

```

report: /Users/lzanda/Documents/REVIEW-TUTORIAL-LLMS-AND-AGENTIC-FOR-SMART-GRIDS/03-codigo-casos-estudio/LLMs-Agents-For-SmartGrids-Code-v2]
↪ /power-flow-agent/benchmarks/results_probe_paper_models/anthropic_claude-opus-4.7/report.json
scoreboard.json: /Users/lzanda/Documents/REVIEW-TUTORIAL-LLMS-AND-AGENTIC-FOR-SMART-GRIDS/03-codigo-casos-estudio/LLMs-Agents-For-SmartGr]
↪ ids-Code-v2/power-flow-agent/benchmarks/results_probe_paper_models/anthropic_claude-opus-4.7/scoreboard.json
scoreboard.csv: /Users/lzanda/Documents/REVIEW-TUTORIAL-LLMS-AND-AGENTIC-FOR-SMART-GRIDS/03-codigo-casos-estudio/LLMs-Agents-For-SmartGri]
↪ ds-Code-v2/power-flow-agent/benchmarks/results_probe_paper_models/anthropic_claude-opus-4.7/scoreboard.csv

```

```

scoreboard.md: /Users/lzanda/Documents/REVIEW-TUTORIAL-LLMS-AND-AGENTIC-FOR-SMART-GRIDS/03-codigo-casos-estudio/LLMs-Agents-For-SmartGrid
  ↳ s-Code-v2/power-flow-agent/benchmarks/results_probe_paper_models/anthropic_claude-opus-4.7/scoreboard.md
scoreboard_per_case.json: /Users/lzanda/Documents/REVIEW-TUTORIAL-LLMS-AND-AGENTIC-FOR-SMART-GRIDS/03-codigo-casos-estudio/LLMs-Agents-Fo
  ↳ r-SmartGrids-Code-v2/power-flow-agent/benchmarks/results_probe_paper_models/anthropic_claude-opus-4.7/scoreboard_per_case.json
traces (40 archivos): /Users/lzanda/Documents/REVIEW-TUTORIAL-LLMS-AND-AGENTIC-FOR-SMART-GRIDS/03-codigo-casos-estudio/LLMs-Agents-For-Sm
  ↳ artGrids-Code-v2/power-flow-agent/benchmarks/results_probe_paper_models/anthropic_claude-opus-4.7/traces
report: /Users/lzanda/Documents/REVIEW-TUTORIAL-LLMS-AND-AGENTIC-FOR-SMART-GRIDS/03-codigo-casos-estudio/LLMs-Agents-For-SmartGrids-Code-v2
  ↳ /power-flow-agent/benchmarks/results_probe_paper_models/openai_gpt-5.4/report.json
scoreboard.json: /Users/lzanda/Documents/REVIEW-TUTORIAL-LLMS-AND-AGENTIC-FOR-SMART-GRIDS/03-codigo-casos-estudio/LLMs-Agents-For-SmartGr
  ↳ ids-Code-v2/power-flow-agent/benchmarks/results_probe_paper_models/openai_gpt-5.4/scoreboard.json
scoreboard.csv: /Users/lzanda/Documents/REVIEW-TUTORIAL-LLMS-AND-AGENTIC-FOR-SMART-GRIDS/03-codigo-casos-estudio/LLMs-Agents-For-SmartGri
  ↳ ds-Code-v2/power-flow-agent/benchmarks/results_probe_paper_models/openai_gpt-5.4/scoreboard.csv
scoreboard.md: /Users/lzanda/Documents/REVIEW-TUTORIAL-LLMS-AND-AGENTIC-FOR-SMART-GRIDS/03-codigo-casos-estudio/LLMs-Agents-For-SmartGrid
  ↳ s-Code-v2/power-flow-agent/benchmarks/results_probe_paper_models/openai_gpt-5.4/scoreboard.md
scoreboard_per_case.json: /Users/lzanda/Documents/REVIEW-TUTORIAL-LLMS-AND-AGENTIC-FOR-SMART-GRIDS/03-codigo-casos-estudio/LLMs-Agents-Fo
  ↳ r-SmartGrids-Code-v2/power-flow-agent/benchmarks/results_probe_paper_models/openai_gpt-5.4/scoreboard_per_case.json
traces (40 archivos): /Users/lzanda/Documents/REVIEW-TUTORIAL-LLMS-AND-AGENTIC-FOR-SMART-GRIDS/03-codigo-casos-estudio/LLMs-Agents-For-Sm
  ↳ artGrids-Code-v2/power-flow-agent/benchmarks/results_probe_paper_models/openai_gpt-5.4/traces
report: /Users/lzanda/Documents/REVIEW-TUTORIAL-LLMS-AND-AGENTIC-FOR-SMART-GRIDS/03-codigo-casos-estudio/LLMs-Agents-For-SmartGrids-Code-v2
  ↳ /power-flow-agent/benchmarks/results_probe_paper_models/openai_gpt-5.5/report.json
scoreboard.json: /Users/lzanda/Documents/REVIEW-TUTORIAL-LLMS-AND-AGENTIC-FOR-SMART-GRIDS/03-codigo-casos-estudio/LLMs-Agents-For-SmartGr
  ↳ ids-Code-v2/power-flow-agent/benchmarks/results_probe_paper_models/openai_gpt-5.5/scoreboard.json
scoreboard.csv: /Users/lzanda/Documents/REVIEW-TUTORIAL-LLMS-AND-AGENTIC-FOR-SMART-GRIDS/03-codigo-casos-estudio/LLMs-Agents-For-SmartGri
  ↳ ds-Code-v2/power-flow-agent/benchmarks/results_probe_paper_models/openai_gpt-5.5/scoreboard.csv
scoreboard.md: /Users/lzanda/Documents/REVIEW-TUTORIAL-LLMS-AND-AGENTIC-FOR-SMART-GRIDS/03-codigo-casos-estudio/LLMs-Agents-For-SmartGrid
  ↳ s-Code-v2/power-flow-agent/benchmarks/results_probe_paper_models/openai_gpt-5.5/scoreboard.md
scoreboard_per_case.json: /Users/lzanda/Documents/REVIEW-TUTORIAL-LLMS-AND-AGENTIC-FOR-SMART-GRIDS/03-codigo-casos-estudio/LLMs-Agents-Fo
  ↳ r-SmartGrids-Code-v2/power-flow-agent/benchmarks/results_probe_paper_models/openai_gpt-5.5/scoreboard_per_case.json
traces (40 archivos): /Users/lzanda/Documents/REVIEW-TUTORIAL-LLMS-AND-AGENTIC-FOR-SMART-GRIDS/03-codigo-casos-estudio/LLMs-Agents-For-Sm
  ↳ artGrids-Code-v2/power-flow-agent/benchmarks/results_probe_paper_models/openai_gpt-5.5/traces
log: /Users/lzanda/Documents/REVIEW-TUTORIAL-LLMS-AND-AGENTIC-FOR-SMART-GRIDS/03-codigo-casos-estudio/LLMs-Agents-For-SmartGrids-Code-v2/po
  ↳ wer-flow-agent/benchmarks/results_probe_paper_models/logs/anthropic_claude-opus-4.7.log
log: /Users/lzanda/Documents/REVIEW-TUTORIAL-LLMS-AND-AGENTIC-FOR-SMART-GRIDS/03-codigo-casos-estudio/LLMs-Agents-For-SmartGrids-Code-v2/po
  ↳ wer-flow-agent/benchmarks/results_probe_paper_models/logs/openai_gpt-5.4.log
log: /Users/lzanda/Documents/REVIEW-TUTORIAL-LLMS-AND-AGENTIC-FOR-SMART-GRIDS/03-codigo-casos-estudio/LLMs-Agents-For-SmartGrids-Code-v2/po
  ↳ wer-flow-agent/benchmarks/results_probe_paper_models/logs/openai_gpt-5.5.log

```