

Informe de experimento: results_stress_gpt-4o-mini_gate3

Generado el 2026-09-15 22:05 por benchmarks/experiment_report.py a partir de 3 report(s) del directorio results_stress_gpt-4o-mini_gate3 (ruta completa en la sección 1). Prosa en español; identificadores (métodos, métricas, campos) en inglés tal como aparecen en los datos.

1. Qué se corrió

Directorio analizado:

```
/Users/lzanda/Documents/REVIEW-TUTORIAL-LLMS-AND-AGENTIC-FOR-SMART-GRIDS/03-codigo-casos-estudio/LLMs-Agents-For-SmartGrids-Code-v2/power-f  
↪ low-agent/benchmarks/results_stress_gpt-4o-mini_gate3
```

- Reports leídos: 3 (report.rescored.json usado en 3; el resto report.json)
- Modelos: openrouter:openai/gpt-4o-mini
- Métodos: pfagent
- Casos: case14, case30, case57
- Requests: generadas, N=40 por caso y seed, difficulties=['stress'] (3 reports)
- Items por dificultad: stress=120
- runs por item: 1; k (perturbación $\pm 10\% \cdot k$): 1; seeds: [0]; max_rounds: 8; temperature: 0.0; timeout_s: 90.0
- Fecha de la corrida (mtime de report.json): 2026-09-14 02:37
- Costo total: 0.2007 USD (120 items; costo medio 0.00167 USD/item)
- Items totales: 120; ok=120; con error=0; tokens totales=1262980
- Solved global: 100.0 % (120/120)

Errores por tipo:

(ningún item con error)

2. Prompts usados

Los prompts se reconstruyen offline con el mismo código del runner (llm.prompt_variants.build_messages sobre la red perturbada con el seed del item). Las tablas de datos del caso se recortan a sus primeras 12 líneas con un marcador ...; el número de caracteres indicado corresponde al prompt completo.

Método pfagent

System prompt (SYSTEM_PROMPT_EN) —1375 caracteres; architecture=react, gate=True, memory=True

You are a power system power-flow analysis assistant. You help the user complete power system analysis tasks from natural-language requests.

```
## What you can do
You complete tasks by calling tool functions:
1. Load an IEEE standard test system (14, 30, 57, 118, 300 buses)
2. Run AC power flow
3. Modify the load at a bus and re-solve
4. Disconnect or reconnect a line and re-solve
5. Query the network status
6. Run an N-1 contingency analysis (disconnect branches one at a time and rank the worst cases)
7. Generate remedial-action suggestions (what-if load shedding or voltage adjustment) to reduce violation risk

## Key rules
- Never fabricate any numerical result. Every voltage, power, loss, or loading value must come from a tool return value.
- If a tool reports an error or a non-converged power flow, say so explicitly and do not report numbers for that state.
- If the request is ambiguous, state the interpretation you are using.
- Bus and line identifiers in the request are MATPOWER 1-based ids; pass them to the tools as given.
- After a network modification, re-solve before reporting any value.
- Always reply in English.

## Reply format
After each analysis, reply in this structure:
1. One-sentence summary
2. Key numbers (total load, total generation, total losses, voltage range) taken from the tool outputs
3. Violations, if any
4. Suggested next step
```

User prompt (primer item) = request_text —90 caracteres; case=case14, seed=1813382119, k=1, request_id=case14-stress-000-s0

Load the IEEE 14-bus system and disconnect the transformer branch between bus 7 and bus 8.

3. Resultados agregados

Columnas del protocolo (medias por método $\times$ caso, del `scoreboard_per_case`): Solved = respondió correctamente de extremo a extremo; Form. = formulación exacta de las tool calls; V_MAE (p.u.) y F_MAE (MW) sobre los items con resultado; FR = tasa de convergencia/feasibility; Faith. = números de la respuesta trazables a salidas de tools; SFR = fallos declarados con seguridad; Claim = éxito reclamado sobre un fallo; Abst. = abstención en items resolubles; Calls = tool calls medias; Tok. = tokens medios; TTV = tiempo medio (s); \$ = costo total USD. Porcentajes en %.

Por método $\times$ caso

Method	Case	N	Solved	Form.	V_MAE	F_MAE	FR	Faith.	SFR	Claim	Abst.	Calls	Tok.	TTV	\$
pfagent	case14	40	100.0	100.0	–	–	0.0	100.0	100.0	0.0	100.0	3.6	6629	3.9	0.0437
pfagent	case30	40	100.0	100.0	0	0.00	2.5	100.0	100.0	0.0	94.1	3.8	13242	4.1	0.0836
pfagent	case57	40	100.0	100.0	–	–	0.0	100.0	100.0	0.0	100.0	2.9	11704	3.7	0.0734

Agregado por método (media ponderada por items sobre todos los casos)

Method	Case	N	Solved	Form.	V_MAE	F_MAE	FR	Faith.	SFR	Claim	Abst.	Calls	Tok.	TTV	\$
pfagent	3 casos	120	100.0	100.0	0	0.00	0.8	100.0	100.0	0.0	40.0	3.4	10525	3.9	0.2007

Formulación exacta por dificultad (items con etapa de tools; n/a = sin tools)

Method	stress
pfagent	100.0 (120/120)

Solved por dificultad (%, items resueltos / items)

Method	stress
pfagent	100.0 (120/120)

Tipos de error de formulación (conteo de items)

Method	ok
pfagent	120

Errores de ejecución por método (conteo de items con error)

(ningún item con *error*)

4. Trazas de ejemplo

Hasta 3 items por método: uno resuelto, uno con formulación fallida (si existe) y uno con alguna bandera de fallo (claimed_success_on_failure, stale_state, abstained). Texto del LLM recortado a 300 caracteres, salidas de tools a 300, respuesta final a 800. Se lee la traza completa de `traces/` cuando existe.

Método pfagent

Ejemplo: resuelto (solved=True)

```
request (case14-stress-000-s0, case14, difficulty=stress, seed=1813382119): Load the IEEE 14-bus system and disconnect the transformer
↳ branch between bus 7 and bus 8.
intended_calls: load_case({"case_name": "case14"}) -> disconnect_line({"from_bus": 7, "to_bus": 8})
executed_calls: load_case({"case_name": "case14"}) -> disconnect_line({"from_bus": 7, "to_bus": 8})
fuente de la traza: archivo de traza completo
--- Ronda 1 ---
tool: load_case({"case_name": "case14"})
  output: {"case_name": "case14", "n_buses": 14, "n_generators": 5, "n_lines": 20, "n_loads": 11, "total_load_mw": 252.6562386075145,
    ↳ "total_gen_capacity_mw": 772.4000000000001}
  gate: n/a (no es un resultado de flujo)
tool: disconnect_line({"from_bus": 7, "to_bus": 8})
```

```

output: {"case_name": "case14", "converged": false, "bus_voltages": [], "line_flows": [], "total_generation_mw": 0.0, "total_load_mw":
  ↳ 0.0, "total_loss_mw": 0.0, "voltage_violations": [], "thermal_violations": [], "summary_text": "", "solver_backend": "pandapower",
  ↳ "llm_prompt": null, "llm_response": null, ...[367 chars más]
gate: FAIL (reasons=isolated_or_unsolved_buses, mismatch=9.78e-09 MW) -> números retenidos

--- Ronda 2 ---
LLM: The IEEE 14-bus system was loaded, and the transformer branch between bus 7 and bus 8 was disconnected. However, the power flow
  ↳ analysis did not converge due to bus 8 being isolated by this topology change.

1. Summary: The network is split, and bus 8 is isolated.
2. Key numbers: Total load and gene ...[293 chars más]
--- Respuesta final ---
The IEEE 14-bus system was loaded, and the transformer branch between bus 7 and bus 8 was disconnected. However, the power flow analysis
  ↳ did not converge due to bus 8 being isolated by this topology change.

1. Summary: The network is split, and bus 8 is isolated.
2. Key numbers: Total load and generation values are not available due to non-convergence.
3. Violations: Bus 8 is isolated, leading to a non-converged solution.
4. Suggested next step: Consider reconnecting the line between bus 7 and bus 8 or modifying the network to restore connectivity. Would you
  ↳ like to reconnect the line?
--- Scoring ---
formulation=ok (exact=True); V_MAE---; F_MAE---; faithful--- %; solved=True (declared_failure); flags=abstained,abstained_on_solvable;
  ↳ error=--; tokens=2039; cost_usd=0.0004; wall=3.6 s

```

```
request (case14-stress-001-s0, case14, difficulty=stress, seed=827308000): Load the 14-bus test case (case14), disconnect the line between
↳ bus 2 and bus 4, disconnect the transformer branch between bus 7 and bus 8, and finally rerun the power flow, then list every line
↳ whose loading is above 100%.
intended_calls: load_case({"case_name": "case14"}) -> disconnect_line({"from_bus": 2, "to_bus": 4}) -> disconnect_line({"from_bus": 7,
↳ "to_bus": 8}) -> run_powerflow({})
executed_calls: load_case({"case_name": "case14"}) -> disconnect_line({"from_bus": 2, "to_bus": 4}) -> disconnect_line({"from_bus": 7,
↳ "to_bus": 8})
fuente de la traza: archivo de traza completo
--- Ronda 1 ---
tool: load_case({"case_name": "case14"})
output: {"case_name": "case14", "n_buses": 14, "n_generators": 5, "n_lines": 20, "n_loads": 11, "total_load_mw": 270.1361769951196,
↳ "total_gen_capacity_mw": 772.4000000000001}
gate: n/a (no es un resultado de flujo)
tool: disconnect_line({"from_bus": 2, "to_bus": 4})
output: {"case_name": "case14", "converged": true, "bus_voltages": [{"bus_id": 1, "vm_pu": 1.06, "va_deg": 0.0, "is_violation": true,
↳ "violation_type": "overvoltage"}, {"bus_id": 2, "vm_pu": 1.0450000000000002, "va_deg": -4.717168458565059, "is_violation": false,
↳ "violation_type": null}, {"bus_id": 3, "vm_p ...[5946 chars más]
gate: PASS (reasons=, mismatch=3.08e-08 MW)
tool: disconnect_line({"from_bus": 7, "to_bus": 8})
output: {"case_name": "case14", "converged": false, "bus_voltages": [], "line_flows": [], "total_generation_mw": 0.0, "total_load_mw":
↳ 0.0, "total_loss_mw": 0.0, "voltage_violations": [], "thermal_violations": [], "summary_text": "", "solver_backend": "pandapower",
↳ "llm_prompt": null, "llm_response": null, ...[367 chars más]
gate: FAIL (reasons=isolated_or_unsolved_buses, mismatch=8.78e-08 MW) -> números retenidos
--- Ronda 2 ---
LLM: The power flow analysis could not be completed due to network isolation caused by the disconnection of the lines.

1. **Summary**: The network is split, and the power flow did not converge.
2. **Key numbers**: Not applicable due to non-convergence.
3. **Violations**: None reported, but the network is ...[163 chars más]
--- Respuesta final ---
The power flow analysis could not be completed due to network isolation caused by the disconnection of the lines.

1. **Summary**: The network is split, and the power flow did not converge.
2. **Key numbers**: Not applicable due to non-convergence.
3. **Violations**: None reported, but the network is isolated.
4. **Suggested next step**: Consider reconnecting one of the disconnected lines to restore network connectivity and then rerun the power
↳ flow analysis.
--- Scoring ---
formulation=ok (exact=True); V_MAE=--; F_MAE=--; faithful=-- %; solved=True (declared_failure); flags=abstained,abstained_on_solvable;
↳ error=--; tokens=4648; cost_usd=0.0008; wall=4.5 s
```

```
request (case14-stress-002-s0, case14, difficulty=stress, seed=1627694679): Load the 14-bus test case (case14), set the active load at bus
↳ 9 to 37.7 MW, disconnect the transformer branch between bus 7 and bus 8, and finally rerun the power flow, then list every line whose
↳ loading is above 80%.
intended_calls: load_case({"case_name": "case14"}) -> modify_load({"bus_id": 9, "p_mw": 37.7}) -> disconnect_line({"from_bus": 7, "to_bus":
↳ 8}) -> run_powerflow({})
executed_calls: load_case({"case_name": "case14"}) -> modify_load({"bus_id": 9, "p_mw": 37.7}) -> disconnect_line({"from_bus": 7, "to_bus":
↳ 8})
fuente de la traza: archivo de traza completo
--- Ronda 1 ---
tool: load_case({"case_name": "case14"})
output: {"case_name": "case14", "n_buses": 14, "n_generators": 5, "n_lines": 20, "n_loads": 11, "total_load_mw": 257.9016499664761,
↳ "total_gen_capacity_mw": 772.4000000000001}
gate: n/a (no es un resultado de flujo)
--- Ronda 2 ---
tool: modify_load({"bus_id": 9, "p_mw": 37.7})
output: {"case_name": "case14", "converged": true, "bus_voltages": [{"bus_id": 1, "vm_pu": 1.06, "va_deg": 0.0, "is_violation": true,
↳ "violation_type": "overvoltage"}, {"bus_id": 2, "vm_pu": 1.0450000000000004, "va_deg": -5.20987075244386, "is_violation": false,
↳ "violation_type": null}, {"bus_id": 3, "vm_pu ...[6276 chars más]
gate: PASS (reasons=-, mismatch=7.12e-09 MW)
```

```

tool: disconnect_line({"from_bus": 7, "to_bus": 8})
  output: {"case_name": "case14", "converged": false, "bus_voltages": [], "line_flows": [], "total_generation_mw": 0.0, "total_load_mw":
    ↪ 0.0, "total_loss_mw": 0.0, "voltage_violations": [], "thermal_violations": [], "summary_text": "", "solver_backend": "pandapower",
    ↪ "llm_prompt": null, "llm_response": null, ...[367 chars más]
  gate: FAIL (reasons=isolated_or_unsolved_buses, mismatch=1.62e-08 MW) -> números retenidos
--- Ronda 3 ---
LLM: The power flow analysis could not be completed due to network isolation caused by the disconnection of the transformer branch between
  ↪ bus 7 and bus 8.

1. **Summary**: The network is split, and the power flow did not converge.
2. **Key numbers**: Not applicable due to non-convergence.
3. **Violation ...[197 chars más]
--- Respuesta final ---
The power flow analysis could not be completed due to network isolation caused by the disconnection of the transformer branch between bus 7
  ↪ and bus 8.

1. **Summary**: The network is split, and the power flow did not converge.
2. **Key numbers**: Not applicable due to non-convergence.
3. **Violations**: None reported due to non-convergence.
4. **Suggested next step**: Consider reconnecting the line between bus 7 and bus 8 to restore network connectivity and then rerun the power
  ↪ flow analysis.
--- Scoring ---
formulation=ok (exact=True); V_MAE=--; F_MAE=--; faithful=-- %; solved=True (declared_failure); flags=abstained,abstained_on_solvable;
  ↪ error=--; tokens=5702; cost_usd=0.0009; wall=4.2 s

```

5. Lecturas

Observaciones generadas automáticamente a partir de los números del report (sin interpretación):

- Todos los métodos que llegan al solver (pfagent) tienen $V_MAE < 1e-4$ p.u. en todos los casos (máximo 0 en pfagent/case30).
- El gate verificó 290 resultados de flujo, 122 fallaron la verificación y se activó (retuvo números) 48 veces en 48 items.
- Ningún item terminó con **error**.
- Solved de **pfagent**: 100.0 % (120/120).
- Formulación exacta agregada sobre métodos con tools: 100.0 % (120/120).
- Ningún item reclamó éxito sobre un fallo (claimed_success_on_failure=0).
- De 119 items en estado de fallo, 119 lo declararon de forma segura (SFR 100.0 % (119/119)).
- 26 de 120 items con mutación de red reportaron estado obsoleto (stale_state).
- Trazabilidad de números (Faith.) media con tools: 100.0 %.
- 119 items esperaban un resultado no convergido (non_converged=71, islanded=48); el método lo declaró correctamente (solved) en 100.0 % (119/119), con solved_reason declared_failure=119.
- Costo total 0.2007 USD; el método más caro fue **pfagent** con 0.2007 USD (100 %).
- Métodos con tools: 3.42 tool calls medias por item (máximo 6).

6. Archivos

Rutas absolutas bajo results_stress_gpt-4o-mini_gate3. report.rescored.json (cuando existe) es el archivo leído para las métricas.

```

report: /Users/lzanda/Documents/REVIEW-TUTORIAL-LLMS-AND-AGENTIC-FOR-SMART-GRIDS/03-codigo-casos-estudio/LLMs-Agents-For-SmartGrids-Code-v2
  ↪ /power-flow-agent/benchmarks/results_stress_gpt-4o-mini_gate3/case14/report.json [leído:
  ↪ report.rescored.json]
scoreboard.json: /Users/lzanda/Documents/REVIEW-TUTORIAL-LLMS-AND-AGENTIC-FOR-SMART-GRIDS/03-codigo-casos-estudio/LLMs-Agents-For-SmartGr
  ↪ ids-Code-v2/power-flow-agent/benchmarks/results_stress_gpt-4o-mini_gate3/case14/scoreboard.json
scoreboard.csv: /Users/lzanda/Documents/REVIEW-TUTORIAL-LLMS-AND-AGENTIC-FOR-SMART-GRIDS/03-codigo-casos-estudio/LLMs-Agents-For-SmartGri
  ↪ ds-Code-v2/power-flow-agent/benchmarks/results_stress_gpt-4o-mini_gate3/case14/scoreboard.csv
scoreboard.md: /Users/lzanda/Documents/REVIEW-TUTORIAL-LLMS-AND-AGENTIC-FOR-SMART-GRIDS/03-codigo-casos-estudio/LLMs-Agents-For-SmartGrid
  ↪ s-Code-v2/power-flow-agent/benchmarks/results_stress_gpt-4o-mini_gate3/case14/scoreboard.md
scoreboard_per_case.json: /Users/lzanda/Documents/REVIEW-TUTORIAL-LLMS-AND-AGENTIC-FOR-SMART-GRIDS/03-codigo-casos-estudio/LLMs-Agents-Fo
  ↪ r-SmartGrids-Code-v2/power-flow-agent/benchmarks/results_stress_gpt-4o-mini_gate3/case14/scoreboard_per_case.json
scoreboard.rescored.json: /Users/lzanda/Documents/REVIEW-TUTORIAL-LLMS-AND-AGENTIC-FOR-SMART-GRIDS/03-codigo-casos-estudio/LLMs-Agents-Fo
  ↪ r-SmartGrids-Code-v2/power-flow-agent/benchmarks/results_stress_gpt-4o-mini_gate3/case14/scoreboard.rescored.json
scoreboard_per_case.rescored.json: /Users/lzanda/Documents/REVIEW-TUTORIAL-LLMS-AND-AGENTIC-FOR-SMART-GRIDS/03-codigo-casos-estudio/LLMs-
  ↪ Agents-For-SmartGrids-Code-v2/power-flow-agent/benchmarks/results_stress_gpt-4o-mini_gate3/case14/scoreboard_per_case.rescored.json
traces (40 archivos): /Users/lzanda/Documents/REVIEW-TUTORIAL-LLMS-AND-AGENTIC-FOR-SMART-GRIDS/03-codigo-casos-estudio/LLMs-Agents-For-Sm
  ↪ artGrids-Code-v2/power-flow-agent/benchmarks/results_stress_gpt-4o-mini_gate3/case14/traces
report: /Users/lzanda/Documents/REVIEW-TUTORIAL-LLMS-AND-AGENTIC-FOR-SMART-GRIDS/03-codigo-casos-estudio/LLMs-Agents-For-SmartGrids-Code-v2
  ↪ /power-flow-agent/benchmarks/results_stress_gpt-4o-mini_gate3/case30/report.json [leído:
  ↪ report.rescored.json]
scoreboard.json: /Users/lzanda/Documents/REVIEW-TUTORIAL-LLMS-AND-AGENTIC-FOR-SMART-GRIDS/03-codigo-casos-estudio/LLMs-Agents-For-SmartGr
  ↪ ids-Code-v2/power-flow-agent/benchmarks/results_stress_gpt-4o-mini_gate3/case30/scoreboard.json
scoreboard.csv: /Users/lzanda/Documents/REVIEW-TUTORIAL-LLMS-AND-AGENTIC-FOR-SMART-GRIDS/03-codigo-casos-estudio/LLMs-Agents-For-SmartGri
  ↪ ds-Code-v2/power-flow-agent/benchmarks/results_stress_gpt-4o-mini_gate3/case30/scoreboard.csv

```

```

scoreboard.md: /Users/lzanda/Documents/REVIEW-TUTORIAL-LLMS-AND-AGENTIC-FOR-SMART-GRIDS/03-codigo-casos-estudio/LLMs-Agents-For-SmartGrid
  ↳ s-Code-v2/power-flow-agent/benchmarks/results_stress_gpt-4o-mini_gate3/case30/scoreboard.md
scoreboard_per_case.json: /Users/lzanda/Documents/REVIEW-TUTORIAL-LLMS-AND-AGENTIC-FOR-SMART-GRIDS/03-codigo-casos-estudio/LLMs-Agents-Fo
  ↳ r-SmartGrids-Code-v2/power-flow-agent/benchmarks/results_stress_gpt-4o-mini_gate3/case30/scoreboard_per_case.json
scoreboard.rescored.json: /Users/lzanda/Documents/REVIEW-TUTORIAL-LLMS-AND-AGENTIC-FOR-SMART-GRIDS/03-codigo-casos-estudio/LLMs-Agents-Fo
  ↳ r-SmartGrids-Code-v2/power-flow-agent/benchmarks/results_stress_gpt-4o-mini_gate3/case30/scoreboard.rescored.json
scoreboard_per_case.rescored.json: /Users/lzanda/Documents/REVIEW-TUTORIAL-LLMS-AND-AGENTIC-FOR-SMART-GRIDS/03-codigo-casos-estudio/LLMs-
  ↳ Agents-For-SmartGrids-Code-v2/power-flow-agent/benchmarks/results_stress_gpt-4o-mini_gate3/case30/scoreboard_per_case.rescored.json
traces (40 archivos): /Users/lzanda/Documents/REVIEW-TUTORIAL-LLMS-AND-AGENTIC-FOR-SMART-GRIDS/03-codigo-casos-estudio/LLMs-Agents-For-Sm
  ↳ artGrids-Code-v2/power-flow-agent/benchmarks/results_stress_gpt-4o-mini_gate3/case30/traces
report: /Users/lzanda/Documents/REVIEW-TUTORIAL-LLMS-AND-AGENTIC-FOR-SMART-GRIDS/03-codigo-casos-estudio/LLMs-Agents-For-SmartGrids-Code-v2
  ↳ /power-flow-agent/benchmarks/results_stress_gpt-4o-mini_gate3/case57/report.json [leído:
  ↳ report.rescored.json]
scoreboard.json: /Users/lzanda/Documents/REVIEW-TUTORIAL-LLMS-AND-AGENTIC-FOR-SMART-GRIDS/03-codigo-casos-estudio/LLMs-Agents-For-SmartGr
  ↳ ids-Code-v2/power-flow-agent/benchmarks/results_stress_gpt-4o-mini_gate3/case57/scoreboard.json
scoreboard.csv: /Users/lzanda/Documents/REVIEW-TUTORIAL-LLMS-AND-AGENTIC-FOR-SMART-GRIDS/03-codigo-casos-estudio/LLMs-Agents-For-SmartGri
  ↳ ds-Code-v2/power-flow-agent/benchmarks/results_stress_gpt-4o-mini_gate3/case57/scoreboard.csv
scoreboard.md: /Users/lzanda/Documents/REVIEW-TUTORIAL-LLMS-AND-AGENTIC-FOR-SMART-GRIDS/03-codigo-casos-estudio/LLMs-Agents-For-SmartGrid
  ↳ s-Code-v2/power-flow-agent/benchmarks/results_stress_gpt-4o-mini_gate3/case57/scoreboard.md
scoreboard_per_case.json: /Users/lzanda/Documents/REVIEW-TUTORIAL-LLMS-AND-AGENTIC-FOR-SMART-GRIDS/03-codigo-casos-estudio/LLMs-Agents-Fo
  ↳ r-SmartGrids-Code-v2/power-flow-agent/benchmarks/results_stress_gpt-4o-mini_gate3/case57/scoreboard_per_case.json
scoreboard.rescored.json: /Users/lzanda/Documents/REVIEW-TUTORIAL-LLMS-AND-AGENTIC-FOR-SMART-GRIDS/03-codigo-casos-estudio/LLMs-Agents-Fo
  ↳ r-SmartGrids-Code-v2/power-flow-agent/benchmarks/results_stress_gpt-4o-mini_gate3/case57/scoreboard.rescored.json
scoreboard_per_case.rescored.json: /Users/lzanda/Documents/REVIEW-TUTORIAL-LLMS-AND-AGENTIC-FOR-SMART-GRIDS/03-codigo-casos-estudio/LLMs-
  ↳ Agents-For-SmartGrids-Code-v2/power-flow-agent/benchmarks/results_stress_gpt-4o-mini_gate3/case57/scoreboard_per_case.rescored.json
traces (40 archivos): /Users/lzanda/Documents/REVIEW-TUTORIAL-LLMS-AND-AGENTIC-FOR-SMART-GRIDS/03-codigo-casos-estudio/LLMs-Agents-For-Sm
  ↳ artGrids-Code-v2/power-flow-agent/benchmarks/results_stress_gpt-4o-mini_gate3/case57/traces
log: /Users/lzanda/Documents/REVIEW-TUTORIAL-LLMS-AND-AGENTIC-FOR-SMART-GRIDS/03-codigo-casos-estudio/LLMs-Agents-For-SmartGrids-Code-v2/po
  ↳ wer-flow-agent/benchmarks/results_stress_gpt-4o-mini_gate3/logs/case14.log
log: /Users/lzanda/Documents/REVIEW-TUTORIAL-LLMS-AND-AGENTIC-FOR-SMART-GRIDS/03-codigo-casos-estudio/LLMs-Agents-For-SmartGrids-Code-v2/po
  ↳ wer-flow-agent/benchmarks/results_stress_gpt-4o-mini_gate3/logs/case30.log
log: /Users/lzanda/Documents/REVIEW-TUTORIAL-LLMS-AND-AGENTIC-FOR-SMART-GRIDS/03-codigo-casos-estudio/LLMs-Agents-For-SmartGrids-Code-v2/po
  ↳ wer-flow-agent/benchmarks/results_stress_gpt-4o-mini_gate3/logs/case57.log

```