

Resumen de bloques: openrouter:openai/gpt-5.4

Generado por benchmarks/build_root_summary.py a partir de 5 bloques. No vuelve a correr nada ni toca los reportes de cada bloque.

Bloques

- results_validation_gpt-5.4_llm_only/REPORT.md -4 filas de scoreboard_per_case
- results_validation_gpt-5.4_agents/REPORT.md -3 filas de scoreboard_per_case
- results_validation_gpt-5.4_prompting/REPORT.md -3 filas de scoreboard_per_case
- results_validation_gpt-5.4_rule_based/REPORT.md -1 filas de scoreboard_per_case
- results_validation_gpt-5.4_stress/REPORT.md -2 filas de scoreboard_per_case

Tabla consolidada -condición normal

Setting	Method	Form.	V_MAE	F_MAE	Solved	Solver status	B_mean	Faith.	SFR	Calls	Tok.	V pass
LLM-only	Structured	-	-	-	0.0 (0/5)	0.0 (0/5)	-	0.0	-	n/a	2390	0.0 (0/5)
LLM-only	Few-shot	-	-	-	-	-	-	-	-	-	-	-
LLM-only	Chain-of-thought	-	-	-	0.0 (0/5)	0.0 (0/5)	-	3.3	-	n/a	2944	0.0 (0/5)
LLM-only	RAG	-	-	-	-	-	-	-	-	-	-	-
Single-call	Structured	20.0 (1/5)	4.33×10^{-3}	1.58	20.0 (1/5)	100.0 (5/5)	4.68×10^{-3}	100.0	-	1.0	4626	100.0 (5/5)
Single-call	Few-shot	20.0 (1/5)	4.33×10^{-3}	1.58	20.0 (1/5)	100.0 (5/5)	4.68×10^{-3}	100.0	-	1.0	4997	100.0 (5/5)
Single-call	Chain-of-thought	20.0 (1/5)	4.33×10^{-3}	1.58	20.0 (1/5)	100.0 (5/5)	4.68×10^{-3}	100.0	-	1.0	5049	100.0 (5/5)
Single-call	RAG	20.0 (1/5)	4.33×10^{-3}	1.58	20.0 (1/5)	100.0 (5/5)	4.68×10^{-3}	100.0	-	1.0	2630	100.0 (5/5)
Architecture	Rule-based parser, no LLM	60.0 (3/5)	0	0.00	60.0 (3/5)	100.0 (3/3)	0	100.0	-	1.6	n/a	100.0 (5/5)
Architecture	Single-call, best prompt	20.0 (1/5)	4.33×10^{-3}	1.58	20.0 (1/5)	100.0 (5/5)	4.68×10^{-3}	100.0	-	1.0	4626	100.0 (5/5)
Architecture	ReAct	100.0 (5/5)	0	0.00	100.0 (5/5)	100.0 (5/5)	0	100.0	-	2.6	9432	100.0 (5/5)
Architecture	Plan-and-Act	-	-	-	-	-	-	-	-	-	-	-
Architecture	ReAct + gate	100.0 (5/5)	0	0.00	100.0 (5/5)	100.0 (5/5)	0	100.0	-	2.6	9443	100.0 (5/5)

Tabla consolidada -condición stress

Setting	Method	Form.	V_MAE	F_MAE	Solved	Solver status	B_mean	Faith.	SFR	Calls	Tok.	V pass
LLM-only	Structured	-	-	-	-	-	-	-	-	-	-	-
LLM-only	Few-shot	-	-	-	-	-	-	-	-	-	-	-
LLM-only	Chain-of-thought	-	-	-	-	-	-	-	-	-	-	-
LLM-only	RAG	-	-	-	-	-	-	-	-	-	-	-
Single-call	Structured	-	-	-	-	-	-	-	-	-	-	-
Single-call	Few-shot	-	-	-	-	-	-	-	-	-	-	-
Single-call	Chain-of-thought	-	-	-	-	-	-	-	-	-	-	-
Single-call	RAG	-	-	-	-	-	-	-	-	-	-	-
Architecture	Rule-based parser, no LLM	-	-	-	-	-	-	-	-	-	-	-
Architecture	Single-call, best prompt	-	-	-	-	-	-	-	-	-	-	-
Architecture	ReAct	50.0 (5/10)	-	0.01	50.0 (5/10)	100.0 (10/10)	0	96.6	100.0 (10/10)	3.6	16333	-
Architecture	Plan-and-Act	-	-	-	-	-	-	-	-	-	-	-
Architecture	ReAct + gate	50.0 (5/10)	-	0.01	50.0 (5/10)	100.0 (10/10)	0	-	100.0 (10/10)	3.9	24422	-