

Informe de experimento: llm_only_structured

Generado el 2026-09-16 17:45 por benchmarks/experiment_report.py a partir de 1 report(s) del directorio llm_only_structured (ruta completa en la sección 1). Prosa en español; identificadores (métodos, métricas, campos) en inglés tal como aparecen en los datos.

1. Qué se corrió

Directorio analizado:

```
/Users/lzanda/Documents/REVIEW-TUTORIAL-LLMS-AND-AGENTIC-FOR-SMART-GRIDS/03-codigo-casos-estudio/LLMs-Agents-For-SmartGrids-Code-v2/power-f ]
↪ low-agent/benchmarks/results_validation_n40/gpt-5.4/llm_only_structured
```

- Reports leídos: 1 (report.rescored.json usado en 1; el resto report.json)
- Modelos: openrouter:openai/gpt-5.4
- Métodos: llm_only:structured
- Casos: case14
- Requests: generadas, N=40 por caso y seed, difficulties=todas (1 report)
- Items por dificultad: plain=10, parameterized=10, multistep=10, ambiguous=10
- runs por item: 1; k (perturbación $\pm 10\% \cdot k$): 1; seeds: [0]; max_rounds: 8; temperature: 0.0; timeout_s: 90.0
- Fecha de la corrida (mtime de report.json): 2026-09-16 13:42
- Costo total: 0.5372 USD (40 items; costo medio 0.01343 USD/item)
- Items totales: 40; ok=40; con error=0; tokens totales=103129
- Solved global: 52.5 % (21/40)

Errores por tipo:

(ningún item con error)

2. Prompts usados

Los prompts se reconstruyen offline con el mismo código del runner (llm.prompt_variants.build_messages sobre la red perturbada con el seed del item). Las tablas de datos del caso se recortan a sus primeras 12 líneas con un marcador ...; el número de caracteres indicado corresponde al prompt completo.

Método llm_only:structured

System prompt —439 caracteres; strategy=structured, forced=False, sin herramientas

You are a power system analysis expert. You need to calculate AC power flow results based on test system data provided by the user.

Important Constraints:

- You must strictly output JSON (no extra text, no Markdown blocks).
- Do not output NaN/Infinity; all values must be finite real numbers.
- If you cannot complete the calculation, set `converged` to false and output the structure you can provide (an empty array is acceptable).

User prompt (primer item) —5127 caracteres; case=case14, seed=1911784258, k=1, request_id=case14-ambiguous-003-s0

System base: 100 MVA. Every per-unit quantity below (r_pu, x_pu, b_pu) is defined on this base. Line and transformer impedances (r_pu, x_pu, b_pu) are already in per unit on this base, not ohms per km; do not convert them further. Branch thermal ratings are not defined for this case, so branch loading is undefined: report loading_percent as 0 for every branch.

You are a power systems expert. Based on the following IEEE case14 test system data, calculate the AC power flow results.

System Data (from pandapower network tables)

Node Data (Bus Table: net.bus)

	name	type	zone	in_service
0	1	b	1.0	True
1	2	b	1.0	True
2	3	b	1.0	True
3	4	b	1.0	True
4	5	b	1.0	True
5	6	b	1.0	True
6	7	b	1.0	True
7	8	b	1.0	True
8	9	b	1.0	True
9	10	b	1.0	True
10	11	b	1.0	True

... (3 líneas omitidas de 15)

Load Data (Load Table: net.load)

	bus	p_mw	q_mvar	in_service
0	2	21.473798	12.567615	True
1	3	85.726702	17.290948	True
2	4	45.568126	-3.717902	True
3	5	6.887322	1.449962	True
4	6	11.740497	7.861940	True
5	9	32.229269	18.135792	True
6	10	9.714580	6.260507	True
7	11	3.424253	1.761044	True
8	12	6.539657	1.715320	True
9	13	14.757610	6.340306	True
10	14	15.783969	5.296634	True

Generator Data (Generator Table: net.gen)

	bus	p_mw	vm_pu	min_p_mw	max_p_mw	min_q_mvar	max_q_mvar	in_service
0	2	41.63323	1.045	0.0	140.0	-40.0	50.0	True
1	3	0.00000	1.010	0.0	100.0	0.0	40.0	True
2	6	0.00000	1.070	0.0	100.0	-6.0	24.0	True
3	8	0.00000	1.090	0.0	100.0	-6.0	24.0	True

Balanced Node/External Grid (External Grid Table: net.ext_grid)

	bus	vm_pu	va_degree	in_service
0	1	1.06	0.0	True

Line Data (Line Table: net.line)

	from_bus	to_bus	r_pu	x_pu	b_pu	parallel	in_service
0	1	2	0.01938	0.05917	0.0528	1	True
1	1	5	0.05403	0.22304	0.0492	1	True
2	2	3	0.04699	0.19797	0.0438	1	True
3	2	4	0.05811	0.17632	0.0340	1	True
4	2	5	0.05695	0.17388	0.0346	1	True
5	3	4	0.06701	0.17103	0.0128	1	True
6	4	5	0.01335	0.04211	0.0000	1	True
7	6	11	0.09498	0.19890	0.0000	1	True
8	6	12	0.12291	0.25581	0.0000	1	True
9	6	13	0.06615	0.13027	0.0000	1	True
10	9	10	0.03181	0.08450	0.0000	1	True

... (4 líneas omitidas de 16)

Transformer Data (Transformer Table: net.trafo)

	hv_bus	lv_bus	r_pu	x_pu	shift_degree	tap_side	tap_neutral	tap_min	tap_max	tap_step_percent	in_service
0	4	7	0.0	0.20912	0.0	hv	0.0	NaN	NaN	2.2	True
1	4	9	0.0	0.55618	0.0	hv	0.0	NaN	NaN	3.1	True
2	5	6	0.0	0.25202	0.0	hv	0.0	NaN	NaN	6.8	True
3	7	8	0.0	0.17615	0.0	None	NaN	NaN	NaN	NaN	True
4	7	9	0.0	0.11001	0.0	None	NaN	NaN	NaN	NaN	True

Task

Load case14 and disconnect the line between bus ten and bus eleven.

Output Requirements

Please calculate and return the following results (JSON format):

- 1) Voltage magnitude (p.u.) and phase angle (degrees) for each node
- 2) Active power (MW) and load factor (%) at the beginning of each branch
 - For lines: `line\_id` uses the row index of `net.line`
 - For transformers: `line\_id` uses the row index of 100000 + `net.trafo`
- 3) Total generation (MW), total load (MW), total losses (MW)
- 4) Are there any voltage overruns (<0.95 or >1.05 p.u.) or line overruns (>100%)?

Please strictly follow the following JSON schema output (all fields are complete and of correct type):

```
{
  "converged": true/false,
  "bus_voltages": [{"bus_id": 1, "vm_pu": 1.0, "va_deg": 0.0}, ...],
  "line_flows": [{"line_id": 0, "p_from_mw": 0.0, "loading_percent": 0.0}, ...],
  "total_generation_mw": 0.0,
  "total_load_mw": 0.0,
  "total_loss_mw": 0.0
}
```


Ejemplo: con bandera de fallo (claimed/stale/abstained)

```
request (case14-ambiguous-031-s0, case14, difficulty=ambiguous, seed=2120122978): Load case14 and disconnect the line between bus six and
↳ bus eleven.
intended_calls: load_case({"case_name": "case14"}) -> disconnect_line({"from_bus": 6, "to_bus": 11})
executed_calls: (sin etapa de tools)
fuente de la traza: archivo de traza completo
(sin rondas de herramientas: una sola llamada al LLM)
--- Respuesta final ---
{"converged":false,"bus_voltages":[],"line_flows":[],"total_generation_mw":0.0,"total_load_mw":0.0,"total_loss_mw":0.0}
--- Scoring ---
formulation=n/a (exact=None); V_MAE--; F_MAE--; faithful=0.0 %; solved=False (convergence_mismatch);
↳ flags=abstained,abstained_on_solvable; error=-; tokens=2052; cost_usd=0.0057; wall=1.3 s
```

Ejemplo: no resuelto

```
request (case14-ambiguous-003-s0, case14, difficulty=ambiguous, seed=1911784258): Load case14 and disconnect the line between bus ten and
↳ bus eleven.
intended_calls: load_case({"case_name": "case14"}) -> disconnect_line({"from_bus": 10, "to_bus": 11})
executed_calls: (sin etapa de tools)
fuente de la traza: archivo de traza completo
(sin rondas de herramientas: una sola llamada al LLM)
--- Respuesta final ---
{"converged":true,"bus_voltages":[{"bus_id":1,"vm_pu":1.06,"va_deg":0.0},{"bus_id":2,"vm_pu":1.045,"va_deg":-4.97},{"bus_id":3,"vm_pu":1.01
,"va_deg":-12.72},{"bus_id":4,"vm_pu":1.017,"va_deg":-10.31},{"bus_id":5,"vm_pu":1.019,"va_deg":-8.77},{"bus_id":6,"vm_pu":1.07,"va_deg
↳ "-14.22},{"bus_id":7,"vm_pu":1.062,"va_deg":-13.37},{"bus_id":8,"vm_pu":1.09,"va_deg":-13.36},{"bus_id":9,"vm_pu":1.056,"va_deg":-14.9
↳ 4},{"bus_id":10,"vm_pu":1.051,"va_deg":-15.1},{"bus_id":11,"vm_pu":1.066,"va_deg":-14.78},{"bus_id":12,"vm_pu":1.055,"va_deg":-15.08},{
↳ "bus_id":13,"vm_pu":1.05,"va_deg":-15.16},{"bus_id":14,"vm_pu":1.034,"va_deg":-16.04}],{"line_id":0,"p_from_mw":156.88,"l
↳ oading_percent":0.0},{"line_id":1,"p_from_mw":75.51,"loading_percent":0.0},{"line_id":2,"p_from_mw":73.24,"loading_perce ...[1022 chars
↳ más]
--- Scoring ---
formulation=n/a (exact=None); V_MAE=2.62e-03; F_MAE=2.10; faithful=0.0 %; solved=False (voltage_error); flags=-; error=-; tokens=2774;
↳ cost_usd=0.0166; wall=6.6 s
```

5. Lecturas

Observaciones generadas automáticamente a partir de los números del report (sin interpretación):

- LLM-only: V_MAE entre 2.05e-03 (llm_only:structured/case14) y 2.05e-03 (llm_only:structured/case14) p.u. sobre los items que devolvieron números.
- 12.5 % (5/40) de abstención en LLM-only (respuesta con converged=false o arrays vacíos).
- En LLM-only, 12 items no resueltos por convergence_mismatch y 7 por voltage_error.
- Ningún item terminó con error.
- Solved de llm_only:structured: 52.5 % (21/40).
- Ningún item reclamó éxito sobre un fallo (claimed_success_on_failure=0).
- Trazabilidad de números (Faith.) media en LLM-only: 1.6 % (sin tools, todo número es no trazable salvo los copiados de la request).
- Costo total 0.5372 USD; el método más caro fue llm_only:structured con 0.5372 USD (100 %).

6. Archivos

Rutas absolutas bajo llm_only_structured.report.rescored.json (cuando existe) es el archivo leído para las métricas.

```
report: /Users/lzanda/Documents/REVIEW-TUTORIAL-LLMS-AND-AGENTIC-FOR-SMART-GRIDS/03-codigo-casos-estudio/LLMs-Agents-For-SmartGrids-Code-v2
↳ /power-flow-agent/benchmarks/results_validation_n40/gpt-5.4/llm_only_structured/report.json [leído:
↳ report.rescored.json]
scoreboard.json: /Users/lzanda/Documents/REVIEW-TUTORIAL-LLMS-AND-AGENTIC-FOR-SMART-GRIDS/03-codigo-casos-estudio/LLMs-Agents-For-SmartGr
↳ ids-Code-v2/power-flow-agent/benchmarks/results_validation_n40/gpt-5.4/llm_only_structured/scoreboard.json
scoreboard.csv: /Users/lzanda/Documents/REVIEW-TUTORIAL-LLMS-AND-AGENTIC-FOR-SMART-GRIDS/03-codigo-casos-estudio/LLMs-Agents-For-SmartGri
↳ ds-Code-v2/power-flow-agent/benchmarks/results_validation_n40/gpt-5.4/llm_only_structured/scoreboard.csv
scoreboard.md: /Users/lzanda/Documents/REVIEW-TUTORIAL-LLMS-AND-AGENTIC-FOR-SMART-GRIDS/03-codigo-casos-estudio/LLMs-Agents-For-SmartGrid
↳ s-Code-v2/power-flow-agent/benchmarks/results_validation_n40/gpt-5.4/llm_only_structured/scoreboard.md
scoreboard_per_case.json: /Users/lzanda/Documents/REVIEW-TUTORIAL-LLMS-AND-AGENTIC-FOR-SMART-GRIDS/03-codigo-casos-estudio/LLMs-Agents-Fo
↳ r-SmartGrids-Code-v2/power-flow-agent/benchmarks/results_validation_n40/gpt-5.4/llm_only_structured/scoreboard_per_case.json
scoreboard.rescored.json: /Users/lzanda/Documents/REVIEW-TUTORIAL-LLMS-AND-AGENTIC-FOR-SMART-GRIDS/03-codigo-casos-estudio/LLMs-Agents-Fo
↳ r-SmartGrids-Code-v2/power-flow-agent/benchmarks/results_validation_n40/gpt-5.4/llm_only_structured/scoreboard.rescored.json
scoreboard_per_case.rescored.json:
↳ /Users/lzanda/Documents/REVIEW-TUTORIAL-LLMS-AND-AGENTIC-FOR-SMART-GRIDS/03-codigo-casos-estudio/LLMs-Agents-For-SmartGrids-Code-v2/p
↳ over-flow-agent/benchmarks/results_validation_n40/gpt-5.4/llm_only_structured/scoreboard_per_case.rescored.json
traces (40 archivos): /Users/lzanda/Documents/REVIEW-TUTORIAL-LLMS-AND-AGENTIC-FOR-SMART-GRIDS/03-codigo-casos-estudio/LLMs-Agents-For-Sm
↳ artGrids-Code-v2/power-flow-agent/benchmarks/results_validation_n40/gpt-5.4/llm_only_structured/traces
```