

Informe de experimento: plan_act_nogate

Generado el 2026-09-19 19:32 por benchmarks/experiment_report.py a partir de 1 report(s) del directorio plan_act_nogate (ruta completa en la sección 1). Prosa en español; identificadores (métodos, métricas, campos) en inglés tal como aparecen en los datos.

1. Qué se corrió

Directorio analizado:

```
/Users/lzanda/Documents/REVIEW-TUTORIAL-LLMS-AND-AGENTIC-FOR-SMART-GRIDS/03-codigo-casos-estudio/LLMs-Agents-For-SmartGrids-Code-v2/power-f  
↪ low-agent/benchmarks/results_validation_plan_structured_controlled/gpt-5.4/plan_act_nogate
```

- Reports leídos: 1 (report.rescored.json usado en 0; el resto report.json)
- Modelos: openrouter:openai/gpt-5.4
- Métodos: plan_act_nogate
- Casos: case14
- Requests: generadas, N=40 por caso y seed, difficulties=todas (1 report)
- Items por dificultad: plain=10, parameterized=10, multistep=10, ambiguous=10
- System prompt: hash 7da5e2a37ec4 (40 filas)
- runs por item: 1; k (perturbación $\pm 10\% \cdot k$): 1; seeds: [0]; max_rounds: 8; temperature: 0.0; timeout_s: 90.0
- Fecha de la corrida (mtime de report.json): 2026-09-19 19:32
- Costo total: 0.0000 USD (40 items; costo medio 0.00000 USD/item)
- Items totales: 40; ok=40; con error=0; tokens totales=127390
- Solved global: 30.0% (12/40)

Errores por tipo:

(ningún item con error)

2. Prompts usados

Los prompts se reconstruyen offline con el mismo código del runner (llm.prompt_variants.build_messages sobre la red perturbada con el seed del item). Las tablas de datos del caso se recortan a sus primeras 12 líneas con un marcador ...; el número de caracteres indicado corresponde al prompt completo. **Esta reconstrucción es ilustrativa, no histórica:** usa el código de llm/prompts.py de HOY, así que si el texto del prompt cambió después de que esta corrida se generó, lo que se ve aquí no es lo que el modelo recibió. El registro histórico real es el hash por fila (system_prompt_hash, ver sección 1 / benchmarks/experiment_report.py::section_what_ran), fijado en el momento de la corrida y nunca recalculado.

Método plan_act_nogate (system_prompt_hash real de esta fila: 7da5e2a37ec4)

System prompt (SYSTEM_PROMPT_EN) (reconstrucción con el código actual, no necesariamente histórica) — 1557 caracteres; architecture=plan_act, gate=False, memory=False

You are a power system power-flow analysis assistant. You help the user complete power system analysis tasks from natural-language requests.

```
## What you can do
You complete tasks by calling tool functions:
1. Load an IEEE standard test system (14, 30, 57, 118, 300 buses)
2. Run AC power flow
3. Modify the load at a bus and re-solve
4. Disconnect or reconnect a line and re-solve
5. Query the network status
6. Run an N-1 contingency analysis (disconnect branches one at a time and rank the worst cases)
7. Generate remedial-action suggestions (what-if load shedding or voltage adjustment) to reduce violation risk

## Key rules
- Never fabricate any numerical result. Every voltage, power, loss, or loading value must come from a tool return value.
- If a tool reports an error or a non-converged power flow, say so explicitly and do not report numbers for that state.
- If the request is ambiguous, state the interpretation you are using.
- Bus and line identifiers in the request are MATPOWER 1-based ids; pass them to the tools as given, unless the request itself states a
  ↪ different indexing convention for that identifier (e.g. "(0-based)"), in which case convert it to the MATPOWER 1-based id before
  ↪ calling any tool.
- After a network modification, re-solve before reporting any value.
- Always reply in English.

## Reply format
After each analysis, reply in this structure:
1. One-sentence summary
2. Key numbers (total load, total generation, total losses, voltage range) taken from the tool outputs
```

- 3. Violations, if any
- 4. Suggested next step

User prompt (primer item) = request_text −67 caracteres; case=case14, seed=1911784258, k=1, request_id=case14-ambiguous-003-s0

Load case14 and disconnect the line between bus ten and bus eleven.

PLAN_SYSTEM_PROMPT (cabecera) −1043 caracteres; primeras 12 líneas de 1424 caracteres; catálogo de tools omitido

You are the planner of a power-system analysis agent. You cannot call tools yourself. Given the user's request, output ONLY a JSON object of the form

```
{
  "plan": [
    {
      "tool": "<tool_name>",
      "args": {
        ...
      }
    },
    ...
  ]
}
```

listing, in execution order, every tool call needed to answer the request. Use only the tools below with exactly these argument names (* = required). Do not include explanations, markdown, or any text outside the JSON.

Available tools:

- load_case(case_name: string*): Load IEEE test case.
- run_powerflow((none)): Run power flow on current case.
- modify_load(bus_id: integer*, p_mw: number*, q_mvar: number): Modify bus load and recompute.
- disconnect_line(from_bus: integer*, to_bus: integer*): Disconnect line/branch between two buses and recompute.
- reconnect_line(from_bus: integer*, to_bus: integer*): Reconnect line/branch between two buses and recompute.
- get_most_loaded_branch((none)): Get most loaded branch.
- run_n1_contingency(top_k: integer, criteria: string, max_candidates: integer): Run N-1 contingency analysis.

...

FINAL_ANSWER_INSTRUCTION (ronda final) −158 caracteres; se añade tras ejecutar el plan

Write the final answer for the user strictly from the tool outputs above. Never invent numbers; if a tool reported an error or a non-converged result, say so.

3. Resultados agregados

Columnas del protocolo (medias por método × caso, del scoreboard_per_case), en los 4 grupos de metricas_pfacent_definicion: **utilidad de tarea** Form. (formulación exacta de las tool calls), V_MAE (p.u., solo sobre peticiones con formulación exacta −toda fila con herramientas que llega al solver con la formulación correcta da exactamente cero, ver benchmarks/scoring.py) y F_MAE (MW) sobre los items con resultado, Solved (respondió correctamente de extremo a extremo); **corrección solver-grounded** Solver status (convergencia reportada = convergencia real), B_mean (residual medio de KCL de los flujos reportados, MW); **fidelidad** Faith. (por respuesta, no por número: fracción de respuestas donde todo número es trazable a una salida de tool Y viene del solve posterior al último cambio de red) y SFR (fallos declarados con seguridad); **costo** Calls (tool calls medias) y Tok. (tokens medios). Fuera de los 4 grupos, solo en este reporte (no en las tablas del paper): Claim (éxito reclamado sobre un fallo), Escalated y Wrong (silent). Escalated y Wrong (silent) junto con solved_autonomously_rate (no la columna Solved de arriba, que puede solaparse con Escalated −una abstención de verificación puede caer sobre un item cuyo cálculo de fondo sí era correcto) suman 100 %: dónde termina cada petición −resuelta sola, entregada a una persona, o mal contestada sin aviso; solo cuenta como Escalated una arquitectura con mecanismo de traspaso real, ver benchmarks.scoring.escalation_check) y \$ (costo total USD). Porcentajes en %.

Por método × caso

Method	Case N	Form.	V_MAE	F_MAE	Solved	Solver status	B_mean	Faith.	SFR	Calls	Tok.	Claim	Escalated	Wrong (silent)	\$
plan_act_nogate	case1440	30.0 (12/40)	0	1.95	30.0 (12/40)	100.0 (40/40)	3.03e-02	95.0 (38/40)	−	1.7	3185−	−	0.0 (0/40)	70.0 (28/40)	−

Formulación exacta por dificultad (items con etapa de tools; n/a = sin tools)

Method	plain	parameterized	multistep	ambiguous
plan_act_nogate	50.0 (5/10)	50.0 (5/10)	0.0 (0/10)	20.0 (2/10)

Solved por dificultad (% , items resueltos / items)

Method	plain	parameterized	multistep	ambiguous
plan_act_nogate	50.0 (5/10)	50.0 (5/10)	0.0 (0/10)	20.0 (2/10)

Tipos de error de formulación (conteo de items)

Method	extra_arg_changes_result	missed_step	ok	wrong_id
plan_act_nogate	6	21	12	1

Errores de ejecución por método (conteo de items con error)

(ningún item con error)

4. Trazas de ejemplo

Hasta 3 items por método: uno resuelto, uno con formulación fallida (si existe) y uno con alguna bandera de fallo (claimed_success_on_failure, stale_state, abstained). Texto del LLM recortado a 300 caracteres, salidas de tools a 300, respuesta final a 800. Se lee la traza completa de `traces/` cuando existe.

Método `plan_act_nogate`

Ejemplo: resuelto (solved=True)

```
request (case14-ambiguous-003-s0, case14, difficulty=ambiguous, seed=1911784258): Load case14 and disconnect the line between bus ten and
↳ bus eleven.
intended_calls: load_case({"case_name": "case14"}) -> disconnect_line({"from_bus": 10, "to_bus": 11})
executed_calls: load_case({"case_name": "case14"}) -> disconnect_line({"from_bus": 10, "to_bus": 11})
fuente de la traza: archivo de traza completo
--- Ronda 1 [plan] ---
tool_call: load_case({"case_name": "case14"})
tool_call: disconnect_line({"from_bus": 10, "to_bus": 11})
--- Ronda 2 [act] ---
tool: load_case({"case_name": "case14"})
output: {"case_name": "case14", "n_buses": 14, "n_generators": 5, "n_lines": 20, "n_loads": 11, "total_load_mw": 253.84578240591338,
↳ "total_gen_capacity_mw": 772.4000000000001}
gate: n/a (no es un resultado de flujo)
tool: disconnect_line({"from_bus": 10, "to_bus": 11})
output: {"case_name": "case14", "converged": true, "bus_voltages": [{"bus_id": 1, "vm_pu": 1.06, "va_deg": 0.0, "is_violation": true,
↳ "violation_type": "overvoltage"}, {"bus_id": 2, "vm_pu": 1.045, "va_deg": -4.782856812869876, "is_violation": false,
↳ "violation_type": null}, {"bus_id": 3, "vm_pu": 1.01, "va_deg": 0.0, "is_violation": false,
↳ "violation_type": null}, {"bus_id": 4, "vm_pu": 1.01, "va_deg": 0.0, "is_violation": false,
↳ "violation_type": null}, {"bus_id": 5, "vm_pu": 1.01, "va_deg": 0.0, "is_violation": false,
↳ "violation_type": null}, {"bus_id": 6, "vm_pu": 1.01, "va_deg": 0.0, "is_violation": false,
↳ "violation_type": null}, {"bus_id": 7, "vm_pu": 1.01, "va_deg": 0.0, "is_violation": false,
↳ "violation_type": null}, {"bus_id": 8, "vm_pu": 1.01, "va_deg": 0.0, "is_violation": false,
↳ "violation_type": null}, {"bus_id": 9, "vm_pu": 1.01, "va_deg": 0.0, "is_violation": false,
↳ "violation_type": null}, {"bus_id": 10, "vm_pu": 1.01, "va_deg": 0.0, "is_violation": false,
↳ "violation_type": null}, {"bus_id": 11, "vm_pu": 1.01, "va_deg": 0.0, "is_violation": false,
↳ "violation_type": null}, {"bus_id": 12, "vm_pu": 1.01, "va_deg": 0.0, "is_violation": false,
↳ "violation_type": null}, {"bus_id": 13, "vm_pu": 1.01, "va_deg": 0.0, "is_violation": false,
↳ "violation_type": null}, {"bus_id": 14, "vm_pu": 1.01, "va_deg": 0.0, "is_violation": false,
↳ "violation_type": null}], "line_flow": [{"line_id": 1, "from_bus": 1, "to_bus": 2, "p_mw": 0.0, "q_mvar": 0.0, "loading": 0.0}, {"line_id": 2, "from_bus": 2, "to_bus": 3, "p_mw": 0.0, "q_mvar": 0.0, "loading": 0.0}, {"line_id": 3, "from_bus": 3, "to_bus": 4, "p_mw": 0.0, "q_mvar": 0.0, "loading": 0.0}, {"line_id": 4, "from_bus": 4, "to_bus": 5, "p_mw": 0.0, "q_mvar": 0.0, "loading": 0.0}, {"line_id": 5, "from_bus": 5, "to_bus": 6, "p_mw": 0.0, "q_mvar": 0.0, "loading": 0.0}, {"line_id": 6, "from_bus": 6, "to_bus": 7, "p_mw": 0.0, "q_mvar": 0.0, "loading": 0.0}, {"line_id": 7, "from_bus": 7, "to_bus": 8, "p_mw": 0.0, "q_mvar": 0.0, "loading": 0.0}, {"line_id": 8, "from_bus": 8, "to_bus": 9, "p_mw": 0.0, "q_mvar": 0.0, "loading": 0.0}, {"line_id": 9, "from_bus": 9, "to_bus": 10, "p_mw": 0.0, "q_mvar": 0.0, "loading": 0.0}, {"line_id": 10, "from_bus": 10, "to_bus": 11, "p_mw": 0.0, "q_mvar": 0.0, "loading": 0.0}, {"line_id": 11, "from_bus": 11, "to_bus": 12, "p_mw": 0.0, "q_mvar": 0.0, "loading": 0.0}, {"line_id": 12, "from_bus": 12, "to_bus": 13, "p_mw": 0.0, "q_mvar": 0.0, "loading": 0.0}, {"line_id": 13, "from_bus": 13, "to_bus": 14, "p_mw": 0.0, "q_mvar": 0.0, "loading": 0.0}], "thermal_violations": []}, {"bus_id": 1, "vm_pu": 1.06, "va_deg": 0.0, "is_violation": true, "violation_type": "overvoltage"}, {"bus_id": 2, "vm_pu": 1.045, "va_deg": -4.782856812869876, "is_violation": false, "violation_type": null}, {"bus_id": 3, "vm_pu": 1.01, "va_deg": 0.0, "is_violation": false, "violation_type": null}, {"bus_id": 4, "vm_pu": 1.01, "va_deg": 0.0, "is_violation": false, "violation_type": null}, {"bus_id": 5, "vm_pu": 1.01, "va_deg": 0.0, "is_violation": false, "violation_type": null}, {"bus_id": 6, "vm_pu": 1.01, "va_deg": 0.0, "is_violation": false, "violation_type": null}, {"bus_id": 7, "vm_pu": 1.01, "va_deg": 0.0, "is_violation": false, "violation_type": null}, {"bus_id": 8, "vm_pu": 1.01, "va_deg": 0.0, "is_violation": false, "violation_type": null}, {"bus_id": 9, "vm_pu": 1.01, "va_deg": 0.0, "is_violation": false, "violation_type": null}, {"bus_id": 10, "vm_pu": 1.01, "va_deg": 0.0, "is_violation": false, "violation_type": null}, {"bus_id": 11, "vm_pu": 1.01, "va_deg": 0.0, "is_violation": false, "violation_type": null}, {"bus_id": 12, "vm_pu": 1.01, "va_deg": 0.0, "is_violation": false, "violation_type": null}, {"bus_id": 13, "vm_pu": 1.01, "va_deg": 0.0, "is_violation": false, "violation_type": null}, {"bus_id": 14, "vm_pu": 1.01, "va_deg": 0.0, "is_violation": false, "violation_type": null}], "line_flow": [{"line_id": 1, "from_bus": 1, "to_bus": 2, "p_mw": 0.0, "q_mvar": 0.0, "loading": 0.0}, {"line_id": 2, "from_bus": 2, "to_bus": 3, "p_mw": 0.0, "q_mvar": 0.0, "loading": 0.0}, {"line_id": 3, "from_bus": 3, "to_bus": 4, "p_mw": 0.0, "q_mvar": 0.0, "loading": 0.0}, {"line_id": 4, "from_bus": 4, "to_bus": 5, "p_mw": 0.0, "q_mvar": 0.0, "loading": 0.0}, {"line_id": 5, "from_bus": 5, "to_bus": 6, "p_mw": 0.0, "q_mvar": 0.0, "loading": 0.0}, {"line_id": 6, "from_bus": 6, "to_bus": 7, "p_mw": 0.0, "q_mvar": 0.0, "loading": 0.0}, {"line_id": 7, "from_bus": 7, "to_bus": 8, "p_mw": 0.0, "q_mvar": 0.0, "loading": 0.0}, {"line_id": 8, "from_bus": 8, "to_bus": 9, "p_mw": 0.0, "q_mvar": 0.0, "loading": 0.0}, {"line_id": 9, "from_bus": 9, "to_bus": 10, "p_mw": 0.0, "q_mvar": 0.0, "loading": 0.0}, {"line_id": 10, "from_bus": 10, "to_bus": 11, "p_mw": 0.0, "q_mvar": 0.0, "loading": 0.0}, {"line_id": 11, "from_bus": 11, "to_bus": 12, "p_mw": 0.0, "q_mvar": 0.0, "loading": 0.0}, {"line_id": 12, "from_bus": 12, "to_bus": 13, "p_mw": 0.0, "q_mvar": 0.0, "loading": 0.0}, {"line_id": 13, "from_bus": 13, "to_bus": 14, "p_mw": 0.0, "q_mvar": 0.0, "loading": 0.0}], "thermal_violations": []}], "gate": PASS (reasons=-, mismatch=7.85e-09 MW)
--- Ronda 3 [answer] [final] ---
LLM: 1. One-sentence summary
```

I loaded IEEE case14 and disconnected the line between bus 10 and bus 11; the post-outage AC power flow converged successfully.

2. Key numbers

```
- Total load: 253.84578240591338 MW
- Total generation: 266.5103144221321 MW
- Total losses: 12.664532024072598 MW
- Voltage range ...[637 chars más]
--- Respuesta final ---
1. One-sentence summary
```

I loaded IEEE case14 and disconnected the line between bus 10 and bus 11; the post-outage AC power flow converged successfully.

2. Key numbers

```
- Total load: 253.84578240591338 MW
- Total generation: 266.5103144221321 MW
- Total losses: 12.664532024072598 MW
- Voltage range: 1.01 pu to 1.09 pu
- Disconnected line flow after outage:
  - Line 10-11: 0.0 MW, 0.0 MVar, 0.0% loading
```

3. Violations, if any

```
- Voltage violations:
  - Bus 1: 1.06 pu, overvoltage
  - Bus 6: 1.0700000000000003 pu, overvoltage
  - Bus 7: 1.0572704509026423 pu, overvoltage
  - Bus 8: 1.09 pu, overvoltage
  - Bus 11: 1.063638182880532 pu, overvoltage
  - Bus 12: 1.0535994495050351 pu, overvoltage
- Thermal violations: none
```

4. Suggested next step

If you want, I can now check how this outage ...[137 chars más]

--- Scoring ---

```
formulation=ok (exact=True); V_MAE=0; F_MAE=0.00; faithful=100.0 %; solved=True (numeric_ok); flags=-; error=-; tokens=3980; cost_usd=-;
↳ wall=4.9 s
```

Ejemplo: formulación fallida (formulation_exact=False)

```
request (case14-ambiguous-011-s0, case14, difficulty=ambiguous, seed=1971893210): Load case14 and set the active load at bus 6 to 14,700 kw.
intended_calls: load_case({"case_name": "case14"}) -> modify_load({"bus_id": 6, "p_mw": 14.7})
executed_calls: load_case({"case_name": "case14"})
fuente de la traza: archivo de traza completo
--- Ronda 1 [plan] ---
tool_call: load_case({"case_name": "case14"})
--- Ronda 2 [act] ---
tool: load_case({"case_name": "case14"})
  output: {"case_name": "case14", "n_buses": 14, "n_generators": 5, "n_lines": 20, "n_loads": 11, "total_load_mw": 263.13580543670497,
    ↪ "total_gen_capacity_mw": 772.4000000000001}
    gate: n/a (no es un resultado de flujo)
--- Ronda 3 [answer] [final] ---
LLM: 1. One-sentence summary
I loaded IEEE case14 successfully, but the active load at bus 6 has not been changed yet because no modification tool output was provided.
```

2. Key numbers

- Case: case14
- Buses: 14
- Generators: 5
- Lines: 20
- Loads: 11
- Total load: 263.13580543670497 MW
- ...[298 chars más]

--- Respuesta final ---

1. One-sentence summary

I loaded IEEE case14 successfully, but the active load at bus 6 has not been changed yet because no modification tool output was provided.

2. Key numbers

- Case: case14
- Buses: 14
- Generators: 5
- Lines: 20
- Loads: 11
- Total load: 263.13580543670497 MW
- Total generation capacity: 772.4000000000001 MW

3. Violations, if any

No power-flow or violation results were provided.

4. Suggested next step

Run the load-modification step to set bus 6 active load to 14,700 kW (14.7 MW), then re-solve the power flow and report the updated system

↪ results.

--- Scoring ---

formulation=missed_step (exact=False); V_MAE=1.54e-04; F_MAE=0.82; faithful=66.7 %; solved=False (formulation); flags=-; error=-;

↪ tokens=1226; cost_usd=-; wall=4.2 s

Ejemplo: no resuelto

```
request (case14-ambiguous-015-s0, case14, difficulty=ambiguous, seed=2078334660): Load the 14-bus test case (case14) and set the active
  ↪ load at bus 5 to 4.4 MVA.
intended_calls: load_case({"case_name": "case14"}) -> modify_load({"bus_id": 5, "p_mw": 4.4})
executed_calls: load_case({"case_name": "case14"})
fuente de la traza: archivo de traza completo
--- Ronda 1 [plan] ---
tool_call: load_case({"case_name": "case14"})
--- Ronda 2 [act] ---
tool: load_case({"case_name": "case14"})
  output: {"case_name": "case14", "n_buses": 14, "n_generators": 5, "n_lines": 20, "n_loads": 11, "total_load_mw": 263.7913959729633,
    ↪ "total_gen_capacity_mw": 772.4000000000001}
    gate: n/a (no es un resultado de flujo)
--- Ronda 3 [answer] [final] ---
LLM: 1. One-sentence summary
The IEEE 14-bus test case (case14) was loaded successfully, but the requested change to set the active load at bus 5 to 4.4 MVA has not
  ↪ been applied in the available tool outputs.
```

2. Key numbers

- Buses: 14
- Generators: 5
- Lines: 20
- Loads: 11
- Total load: 26 ...[323 chars más]

--- Respuesta final ---

1. One-sentence summary

The IEEE 14-bus test case (case14) was loaded successfully, but the requested change to set the active load at bus 5 to 4.4 MVA has not
↪ been applied in the available tool outputs.

2. Key numbers

- Buses: 14
- Generators: 5
- Lines: 20
- Loads: 11
- Total load: 263.7913959729633 MW
- Total generation capacity: 772.4000000000001 MW

3. Violations, if any

No power-flow results or violation results are available from the tool outputs provided.

4. Suggested next step

Apply the load modification at bus 5 and then run a power flow so the updated operating point can be reported.

--- Scoring ---

```
formulation=missed_step (exact=False); V_MAE=7.51e-05; F_MAE=0.28; faithful=100.0 %; solved=False (formulation); flags=-; error=-;
↳ tokens=1247; cost_usd=-; wall=3.9 s
```

5. Lecturas

Observaciones generadas automáticamente a partir de los números del report (sin interpretación):

- V_MAE $\geq 1e-4$ p.u. en 1 filas método×caso con solver: plan_act_nogate/case14 (1.63e-03).
- El gate verificó 22 resultados de flujo, 0 fallaron la verificación y se activó (retuvo números) 0 veces en 0 items.
- Ningún item terminó con **error**.
- Solved de plan_act_nogate: 30.0 % (12/40).
- Formulación exacta agregada sobre métodos con tools: 30.0 % (12/40).
- Formulación por dificultad: mejor plain 50.0 % (5/10), peor multistep 0.0 % (0/10).
- Errores de formulación más frecuentes: missed_step=21, extra_arg_changes_result=6, wrong_id=1.
- Ningún item reclamó éxito sobre un fallo (claimed_success_on_failure=0).
- 0 de 11 items con mutación de red reportaron estado obsoleto (stale_state).
- Trazabilidad de números (Faith.) media con tools: 99.0 %.
- Métodos con tools: 1.73 tool calls medias por item (máximo 6).

6. Archivos

Rutas absolutas bajo plan_act_nogate. report.rescored.json (cuando existe) es el archivo leído para las métricas.

```
report: /Users/lzanda/Documents/REVIEW-TUTORIAL-LLMS-AND-AGENTIC-FOR-SMART-GRIDS/03-codigo-casos-estudio/LLMs-Agents-For-SmartGrids-Code-v2/
↳ /power-flow-agent/benchmarks/results_validation_plan_structured_controlled/gpt-5.4/plan_act_nogate/report.json
scoreboard.json: /Users/lzanda/Documents/REVIEW-TUTORIAL-LLMS-AND-AGENTIC-FOR-SMART-GRIDS/03-codigo-casos-estudio/LLMs-Agents-For-SmartGrids-Code-v2/power-flow-agent/benchmarks/results_validation_plan_structured_controlled/gpt-5.4/plan_act_nogate/scoreboard.json
scoreboard.csv: /Users/lzanda/Documents/REVIEW-TUTORIAL-LLMS-AND-AGENTIC-FOR-SMART-GRIDS/03-codigo-casos-estudio/LLMs-Agents-For-SmartGrids-Code-v2/power-flow-agent/benchmarks/results_validation_plan_structured_controlled/gpt-5.4/plan_act_nogate/scoreboard.csv
scoreboard.md: /Users/lzanda/Documents/REVIEW-TUTORIAL-LLMS-AND-AGENTIC-FOR-SMART-GRIDS/03-codigo-casos-estudio/LLMs-Agents-For-SmartGrids-Code-v2/power-flow-agent/benchmarks/results_validation_plan_structured_controlled/gpt-5.4/plan_act_nogate/scoreboard.md
scoreboard_per_case.json:
↳ /Users/lzanda/Documents/REVIEW-TUTORIAL-LLMS-AND-AGENTIC-FOR-SMART-GRIDS/03-codigo-casos-estudio/LLMs-Agents-For-SmartGrids-Code-v2/power-flow-agent/benchmarks/results_validation_plan_structured_controlled/gpt-5.4/plan_act_nogate/scoreboard_per_case.json
traces (40 archivos): /Users/lzanda/Documents/REVIEW-TUTORIAL-LLMS-AND-AGENTIC-FOR-SMART-GRIDS/03-codigo-casos-estudio/LLMs-Agents-For-SmartGrids-Code-v2/power-flow-agent/benchmarks/results_validation_plan_structured_controlled/gpt-5.4/plan_act_nogate/traces
```